

---

# Announcements

---

- ❖ Mid-term exam: June 22, 4pm-5:40pm
  - ❖ Open book, open notes
  - ❖ No communication
- ❖ HW5 on CSP
  - ❖ Released today
  - ❖ Due June 24 at 11:59pm
- ❖ P3 on MDP and RL
  - ❖ Early release
  - ❖ Due July 3 at 11:59pm

# Ve492: Introduction to Artificial Intelligence

## Mid-term Review

---



Paul Weng

UM-SJTU Joint Institute

Slides adapted from <http://ai.berkeley.edu>, AIMA, UM, CMU

# What have we learned so far?

- ❖ Search and planning

- ❖ Define a state space, goal test; Find path from start to goal

- ❖ Game trees

- ❖ Define utilities; Find path from start that maximizes utility

- ❖ Decision theory and game theory

- ❖ Foundation for MEU; Basic concepts in game theory

- ❖ MDPs

- ❖ Define rewards, utility = (discounted) sum of rewards
- ❖ Find policy that maximizes utility

- ❖ Reinforcement learning

- ❖ Just like MDPs, only T and / or R are not known in advance



- ❖ Constraint satisfaction

- ❖ Find solution that satisfies constraints; Not just for finding a sequential plan



# High-Level Framework

- ❖ How to build AI system?



---

# Search

---

- ❖ Environment: single-agent, fully-observable state, deterministic transition, sequential, model known
- ❖ Search problem
  - ❖ States, transition model, goal test, initial state
  - ❖ Search tree
- ❖ Algorithms
  - ❖ Uninformed search
    - ❖ BFS, DFS, UCS
  - ❖ Informed search
    - ❖ Greedy search,  $A^*$
- ❖ Properties
  - ❖ Complete, optimal
  - ❖ Space and computational complexities

# Search in Games

---

- ❖ Environment: multi-agent, fully-observable state, deterministic or stochastic transition, turn-taking , model known
- ❖ Multi-agent search problems as games
  - ❖ States, players, transition model, terminal test/ values, initial state
  - ❖ Game tree
- ❖ Algorithm for adversarial agent (zero-sum game)
  - ❖ Minimax search algorithm
  - ❖ Alpha-beta pruning
  - ❖ Depth-limited search, iterative deepening
- ❖ Algorithm for random agent
  - ❖ Expectimax
- ❖ Algorithm for multi-agent search
  - ❖ Expectiminimax

# Decision Theory and Game Theory

---

## ❖ Axiomatization of Expected Utility

- ❖ Completeness, Transitivity, Independence, Continuity
- ❖ Unicity of utility function up to positive affine transformation
- ❖ Preference elicitation

## ❖ Game theory

- ❖ Extensive form vs normal form
- ❖ Best response, dominant/ dominated strategies
- ❖ Nash equilibrium (pure or mixed)
- ❖ Pareto optimal, correlated equilibrium

# Markov Decision Process

---

- ❖ Environment: single-agent, fully-observable state, stochastic transition, sequential, model known
- ❖ Model
  - ❖ States, actions, transition function, reward function
- ❖ Algorithms
  - ❖ Policy evaluation
  - ❖ Policy extraction
  - ❖ Value iteration
  - ❖ Policy iteration



# Reinforcement Learning

---

- ❖ Environment: single-agent, fully-observable state, stochastic transition, sequential, model unknown
- ❖ MDP Model, but unknown!
  - ❖ States, actions, transition function, reward function
- ❖ Algorithms
  - ❖ Policy evaluation with TD learning
  - ❖ Policy learning with Q-learning
  - ❖ Approximate Q-learning
  - ❖ Action selection with  $\epsilon$ -greedy or exploration function

# Constraint Satisfaction

---

- ❖ CSP

- ❖ Set of variables, set of domains, set of constraints
- ❖ Find assignments to variables such that all constraints are satisfied

- ❖ Algorithms

- ❖ Backtracking search
  - ❖ Filtering, forward-checking, arc consistency, k-consistency
  - ❖ Ordering of variables and values
- ❖ Structure of constraint graph
  - ❖ Two-pass algorithm for tree-structured constraint graph
  - ❖ Cutset conditioning
- ❖ Iterative improvement

- ❖ Local search

# Quiz: Search

- ❖ Consider a graph search problem where for every action, the cost is at least  $\epsilon$ , with  $\epsilon > 0$ . Assume the used heuristic is consistent.
- ❖ Greedy graph search is guaranteed to return an optimal solution. **F**
- ❖ A\* graph search is guaranteed to return an optimal solution. **T**
- ❖ A\* graph search is guaranteed to expand no more nodes than depth-first graph search. **F**
- ❖ A\* graph search is guaranteed to expand no more nodes than uniform-cost graph search. **T**

# Quiz: A\* Heuristics

---

- ❖ Let  $H_1$  and  $H_2$  both be admissible heuristics.
  - ❖  $\max(H_1, H_2)$  is necessarily admissible T
  - ❖  $\min(H_1, H_2)$  is necessarily admissible T
  - ❖  $(H_1 + H_2)/2$  is necessarily admissible T
  - ❖  $\max(H_1, H_2)$  is necessarily consistent F

# Quiz: Search under Uncertainty

- ❖ You are given a game tree for which you are the maximizer, and in the nodes in which you don't get to make a decision an action is chosen uniformly at random amongst the available options. Your objective is to maximize the probability you win \$10 or more (rather than the usual objective to maximize your expected value).
- ❖ Running expectimax will result in finding the optimal strategy to maximize the probability of winning \$10 or more. **F**
- ❖ Running minimax, where chance nodes are considered minimizers, will result in finding the optimal strategy to maximize the probability of winning \$10 or more. **F**
- ❖ Running expectimax in a modified game tree where every pay-off of \$10 or more is given a value of 1, and every pay-off lower than \$10 is given a value of 0 will result in finding the optimal strategy to maximize the probability of winning \$10 or more. **T**
- ❖ Running minimax in a modified game tree where every pay-off of \$10 or more is given a value of 1, and every pay-off lower than \$10 is given a value of 0 will result in finding the optimal strategy to maximize the probability of winning \$10 or more. **F**

# Quiz: Adversarial Search

- ❖ In the context of adversarial search,  $\alpha$ - $\beta$  pruning
  - ❖ can reduce computation time by pruning portions of the game tree **T**
  - ❖ is generally faster than minimax, but loses the guarantee of optimality **F**
  - ❖ always returns the same value as minimax for the root of the tree **T**
  - ❖ always returns the same value as minimax for all nodes of the tree **F**

# Game Theory: Zero-Sum Game

---

- ❖ Two players choose simultaneously a coin of 10 cents, 50 cents or 1 dollar, which they show to each other.
- ❖ If they chose the same coin, player I wins. Otherwise, player II wins.
- ❖ Write this game in normal form. Is there any pure NE? No
- ❖ Express a system of inequalities to find a mixed NE.

# Quiz: MDP

- ❖ For Markov Decisions Processes (MDPs), we have that:
  - ❖ A small discount (close to 0) encourages shortsighted, greedy behavior. **T**
  - ❖ A large, negative **living reward** ( $\ll 0$ ) encourages shortsighted, greedy behavior. **T** meaning every time step will add a negative reward
  - ❖ A negative living reward can always be expressed using a discount  $< 1$ . **F** There is no direct relation between these two concepts.
  - ❖ A discount  $< 1$  can always be expressed as a negative living reward. **F**



# Quiz: MDP

- ❖ Value iteration can converge only if the discount factor ( $\gamma$ ) satisfies  $0 < \gamma < 1$ . F It can sometimes converge even if discount factor is equal to 1.
- ❖ Policies found by value iteration may be superior to policies found by policy iteration. F
- ❖ Policies found by policy iteration may be superior to policies found by value iteration. F The policy found by these two method should be the same.
- ❖ In some problems, value iteration can converge even though policy iteration may not. T not very important if you don't understand


# Quiz: Reinforcement Learning

- ❖ Assume that the agent observes the true reward with some Gaussian noise  $\mathcal{N}(0,1)$ , Q-learning would still converge. **T**  
Because Q-learning do the average, and even there is noise in the average, it will still work.
- ❖ Q-learning can learn the optimal Q-function  $Q^*$  without ever executing the optimal policy. **T**  
Becuase Q-learning is off-policy.
- ❖ If an MDP has a transition model  $T$  that assigns non-zero probability for all triples  $T(s, a, s')$  then Q-learning will fail. **F**
- ❖ In Q-learning, we decide to explore every  $k$  steps, i.e., if  $t = 0 [k]$  we choose a random action with a uniform distribution, otherwise we choose the greedy action. This version would still converge. **F**

# Quiz: CSP

- ❖ Assume given a CSP whose constraint graph is given below and that all the variables have the same domain.  
29 variables; d size of domain
- ❖ What is the complexity of solving it with a direct application of backtracking search?  $O(d^{29})$
- ❖ Which efficient strategy could you apply to solve it? What would be the complexity?

$$O(d \cdot d^2 \cdot 28) = O(d^3)$$



# CSP Problem: Job Scheduling

## ❖ When can I move in?

| Task | Description        | Duration | Predecessor |
|------|--------------------|----------|-------------|
| a    | Erecting walls     | 7        | none        |
| b    | Carpentry for roof | 3        | a           |
| c    | Roof               | 1        | b           |
| d    | Installations      | 8        | a           |
| e    | Facade painting    | 2        | c & d       |
| f    | Windows            | 1        | c & d       |
| g    | Garden             | 1        | c & d       |
| h    | Ceilings           | 3        | a           |
| i    | Painting           | 2        | f & h       |
| j    | Moving in          | 1        | i           |