

Ve492: Introduction to Artificial Intelligence

Classical Planning

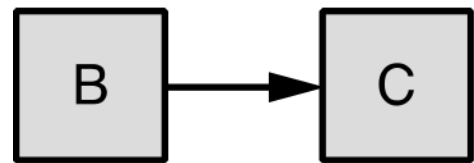


Paul Weng

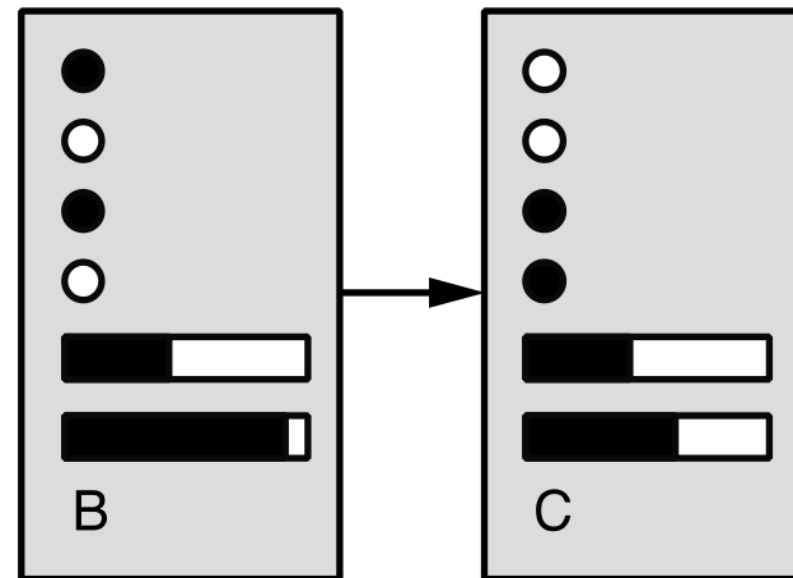
UM-SJTU Joint Institute

Slides adapted from CMU, AIMA, UM

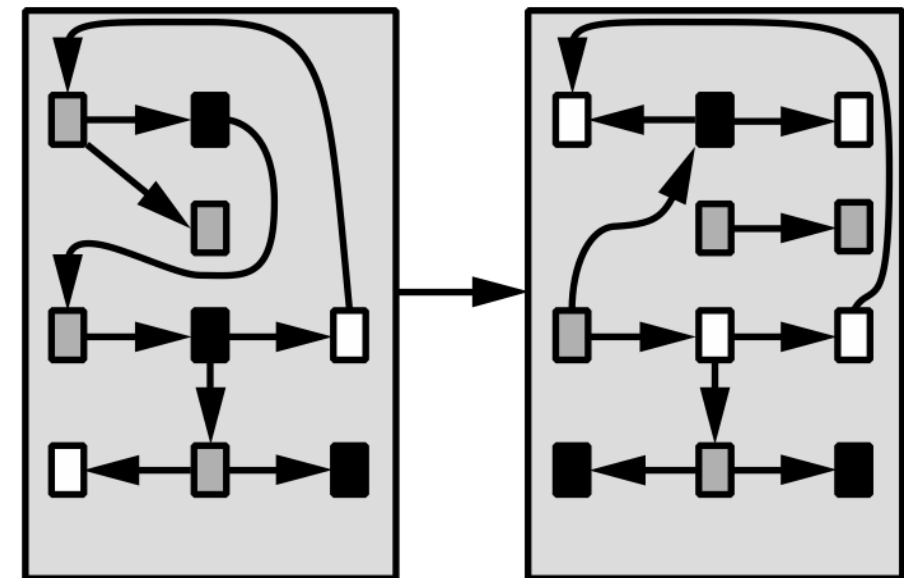
Spectrum of representations



(a) Atomic



(b) Factored



(b) Structured

**Search,
game-playing
MDP, RL**

**CSPs, planning,
propositional logic,
Bayes nets, neural
nets, RL with
function approx.**

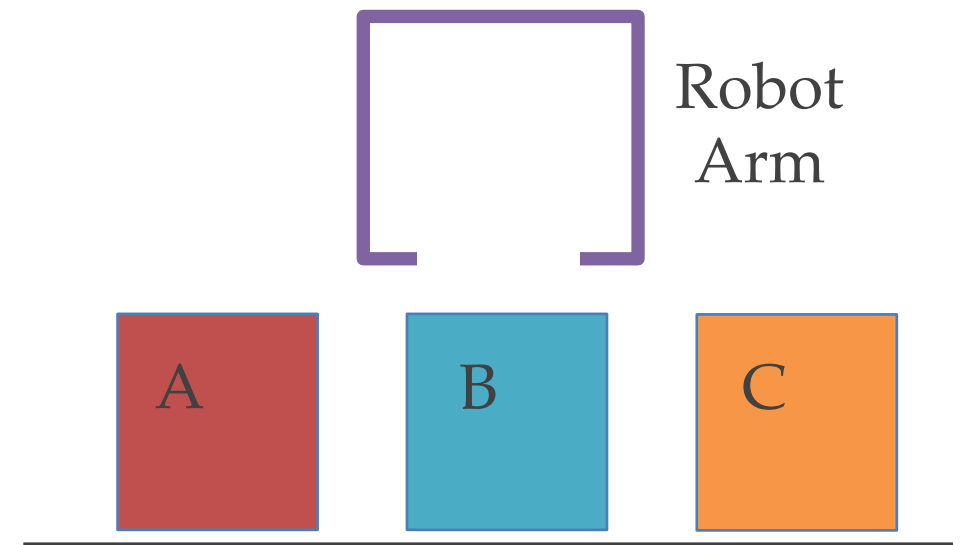
**First-order logic,
databases,
probabilistic programs**

Robot Block Stacking



Start state: A, B, C on table
Goal: Block B on C and C on A
Plan: ?

Modeling Block Stacking States



Start state: A, B, C on table
Goal: Block B on C and C on A
Plan: ?

Goals in the World

Goal States
completely specified

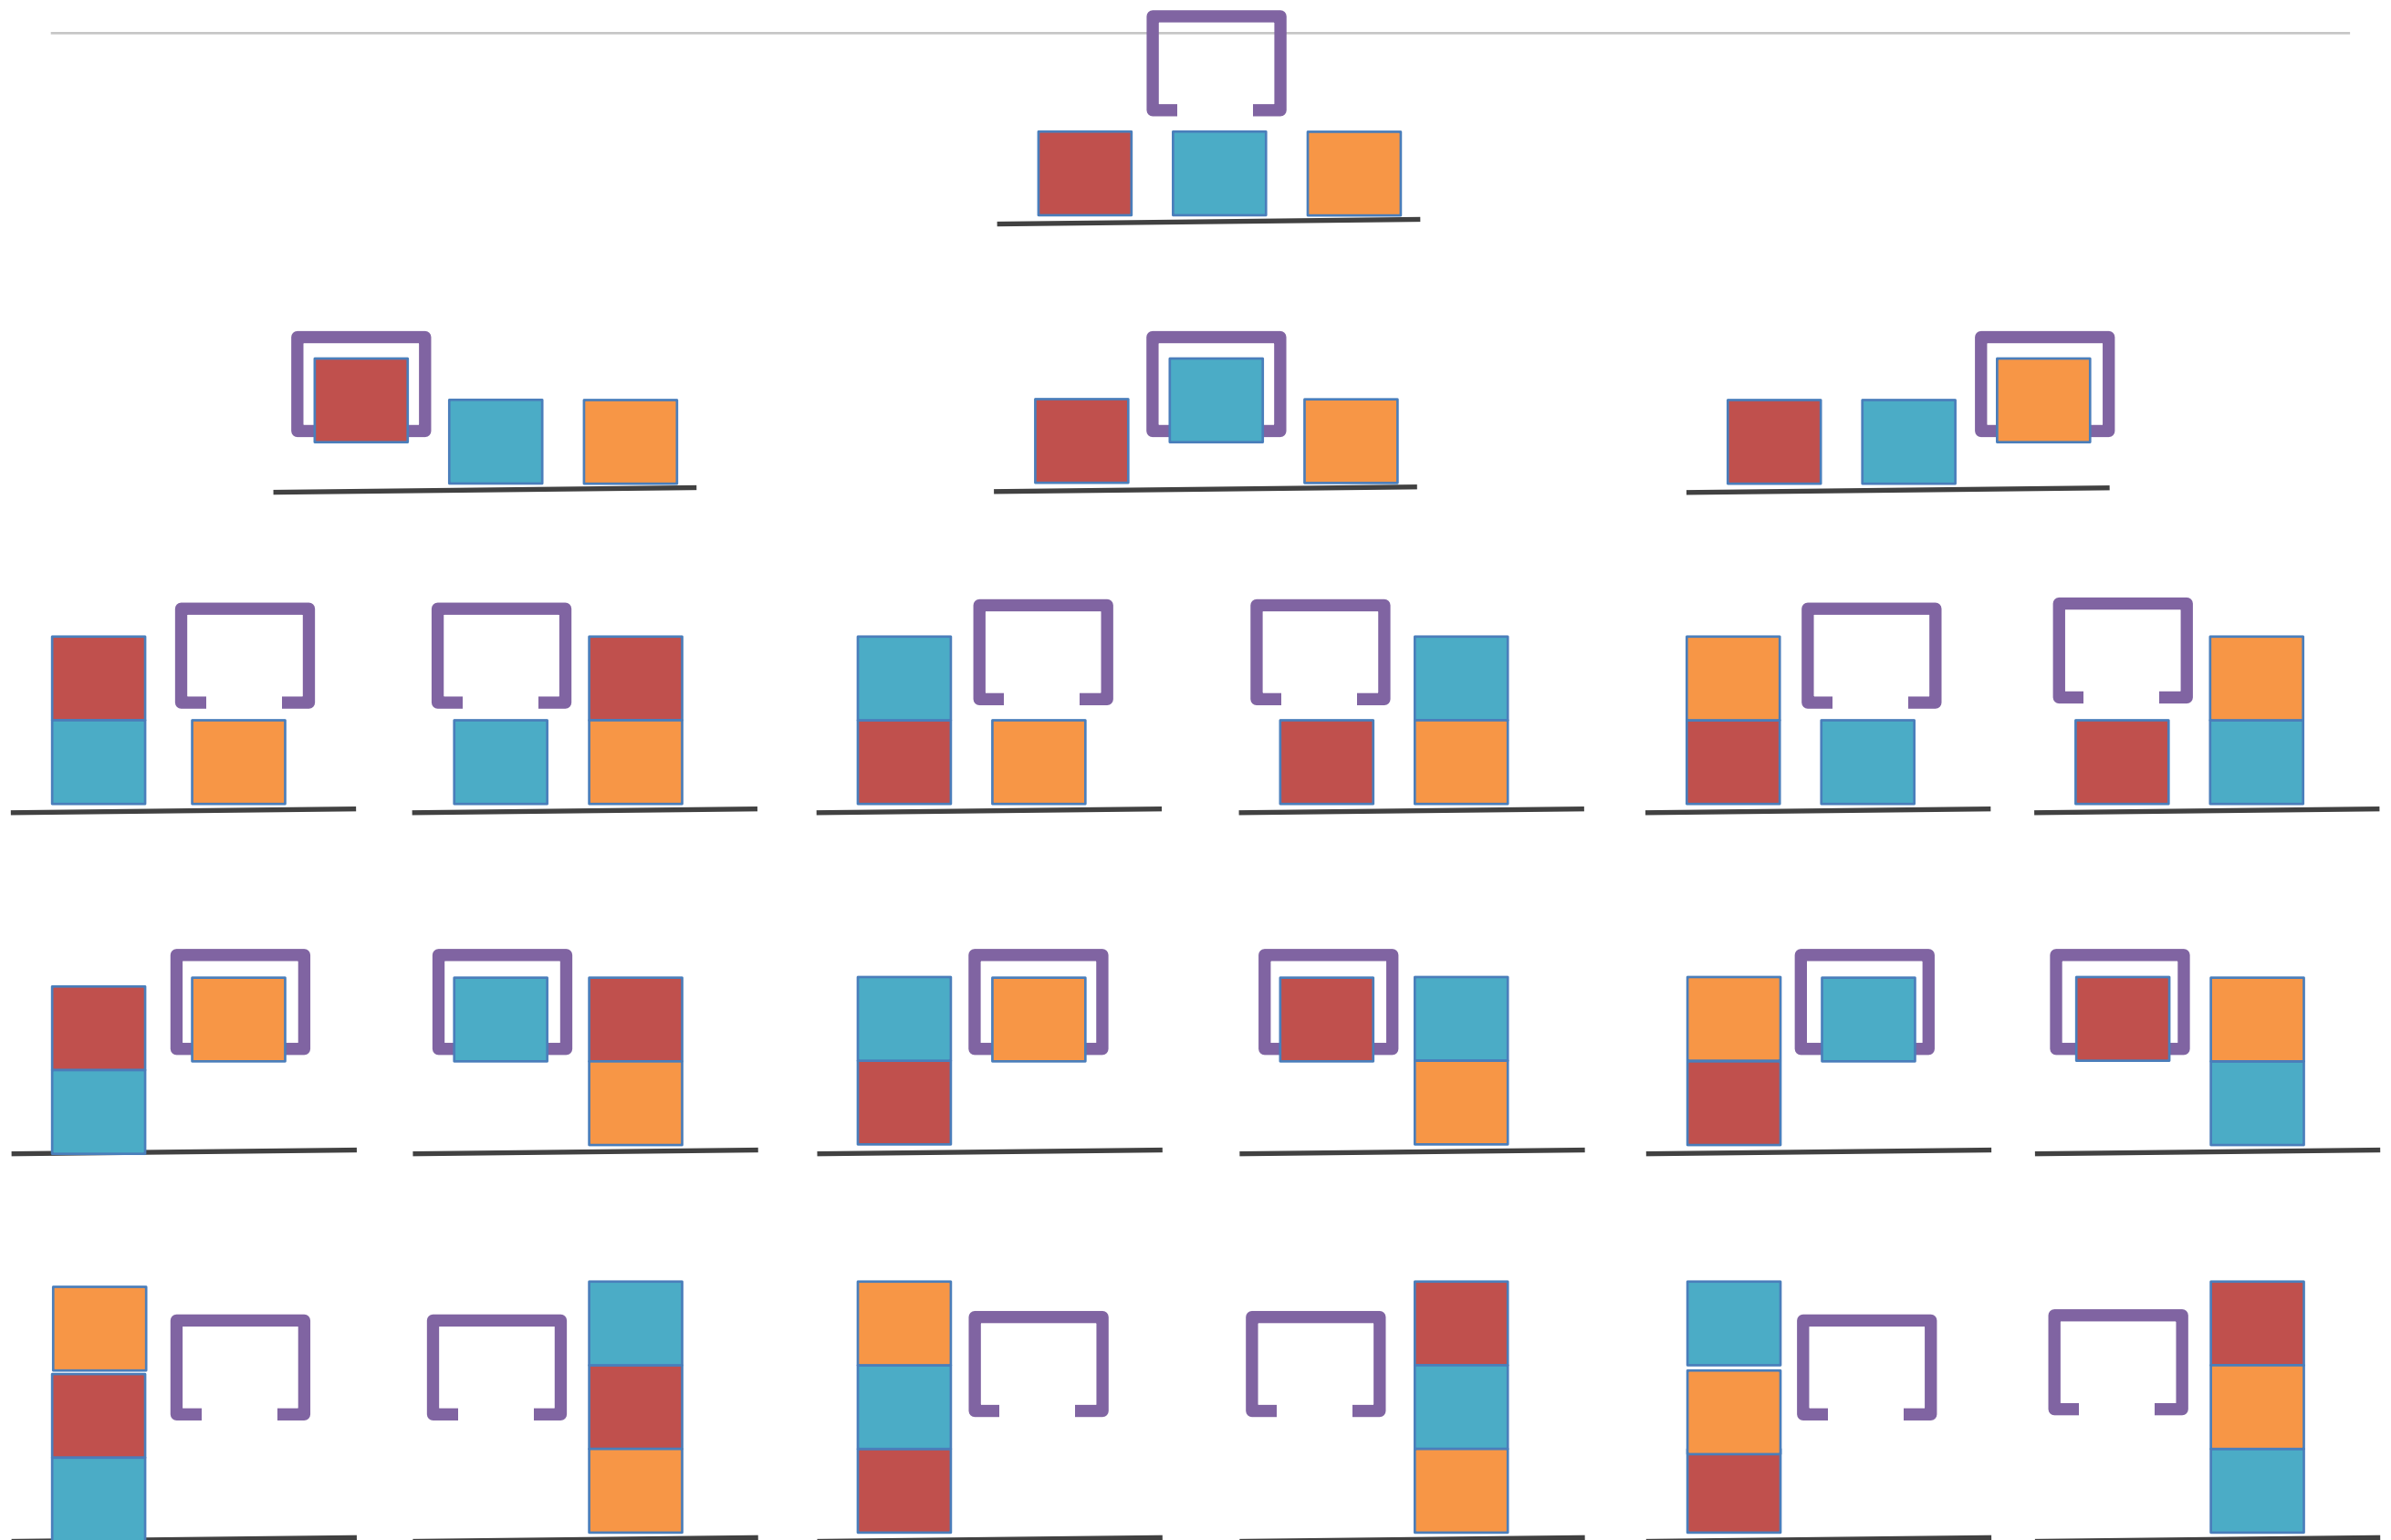
Goal Statements
partially specified

Preference models
objective function

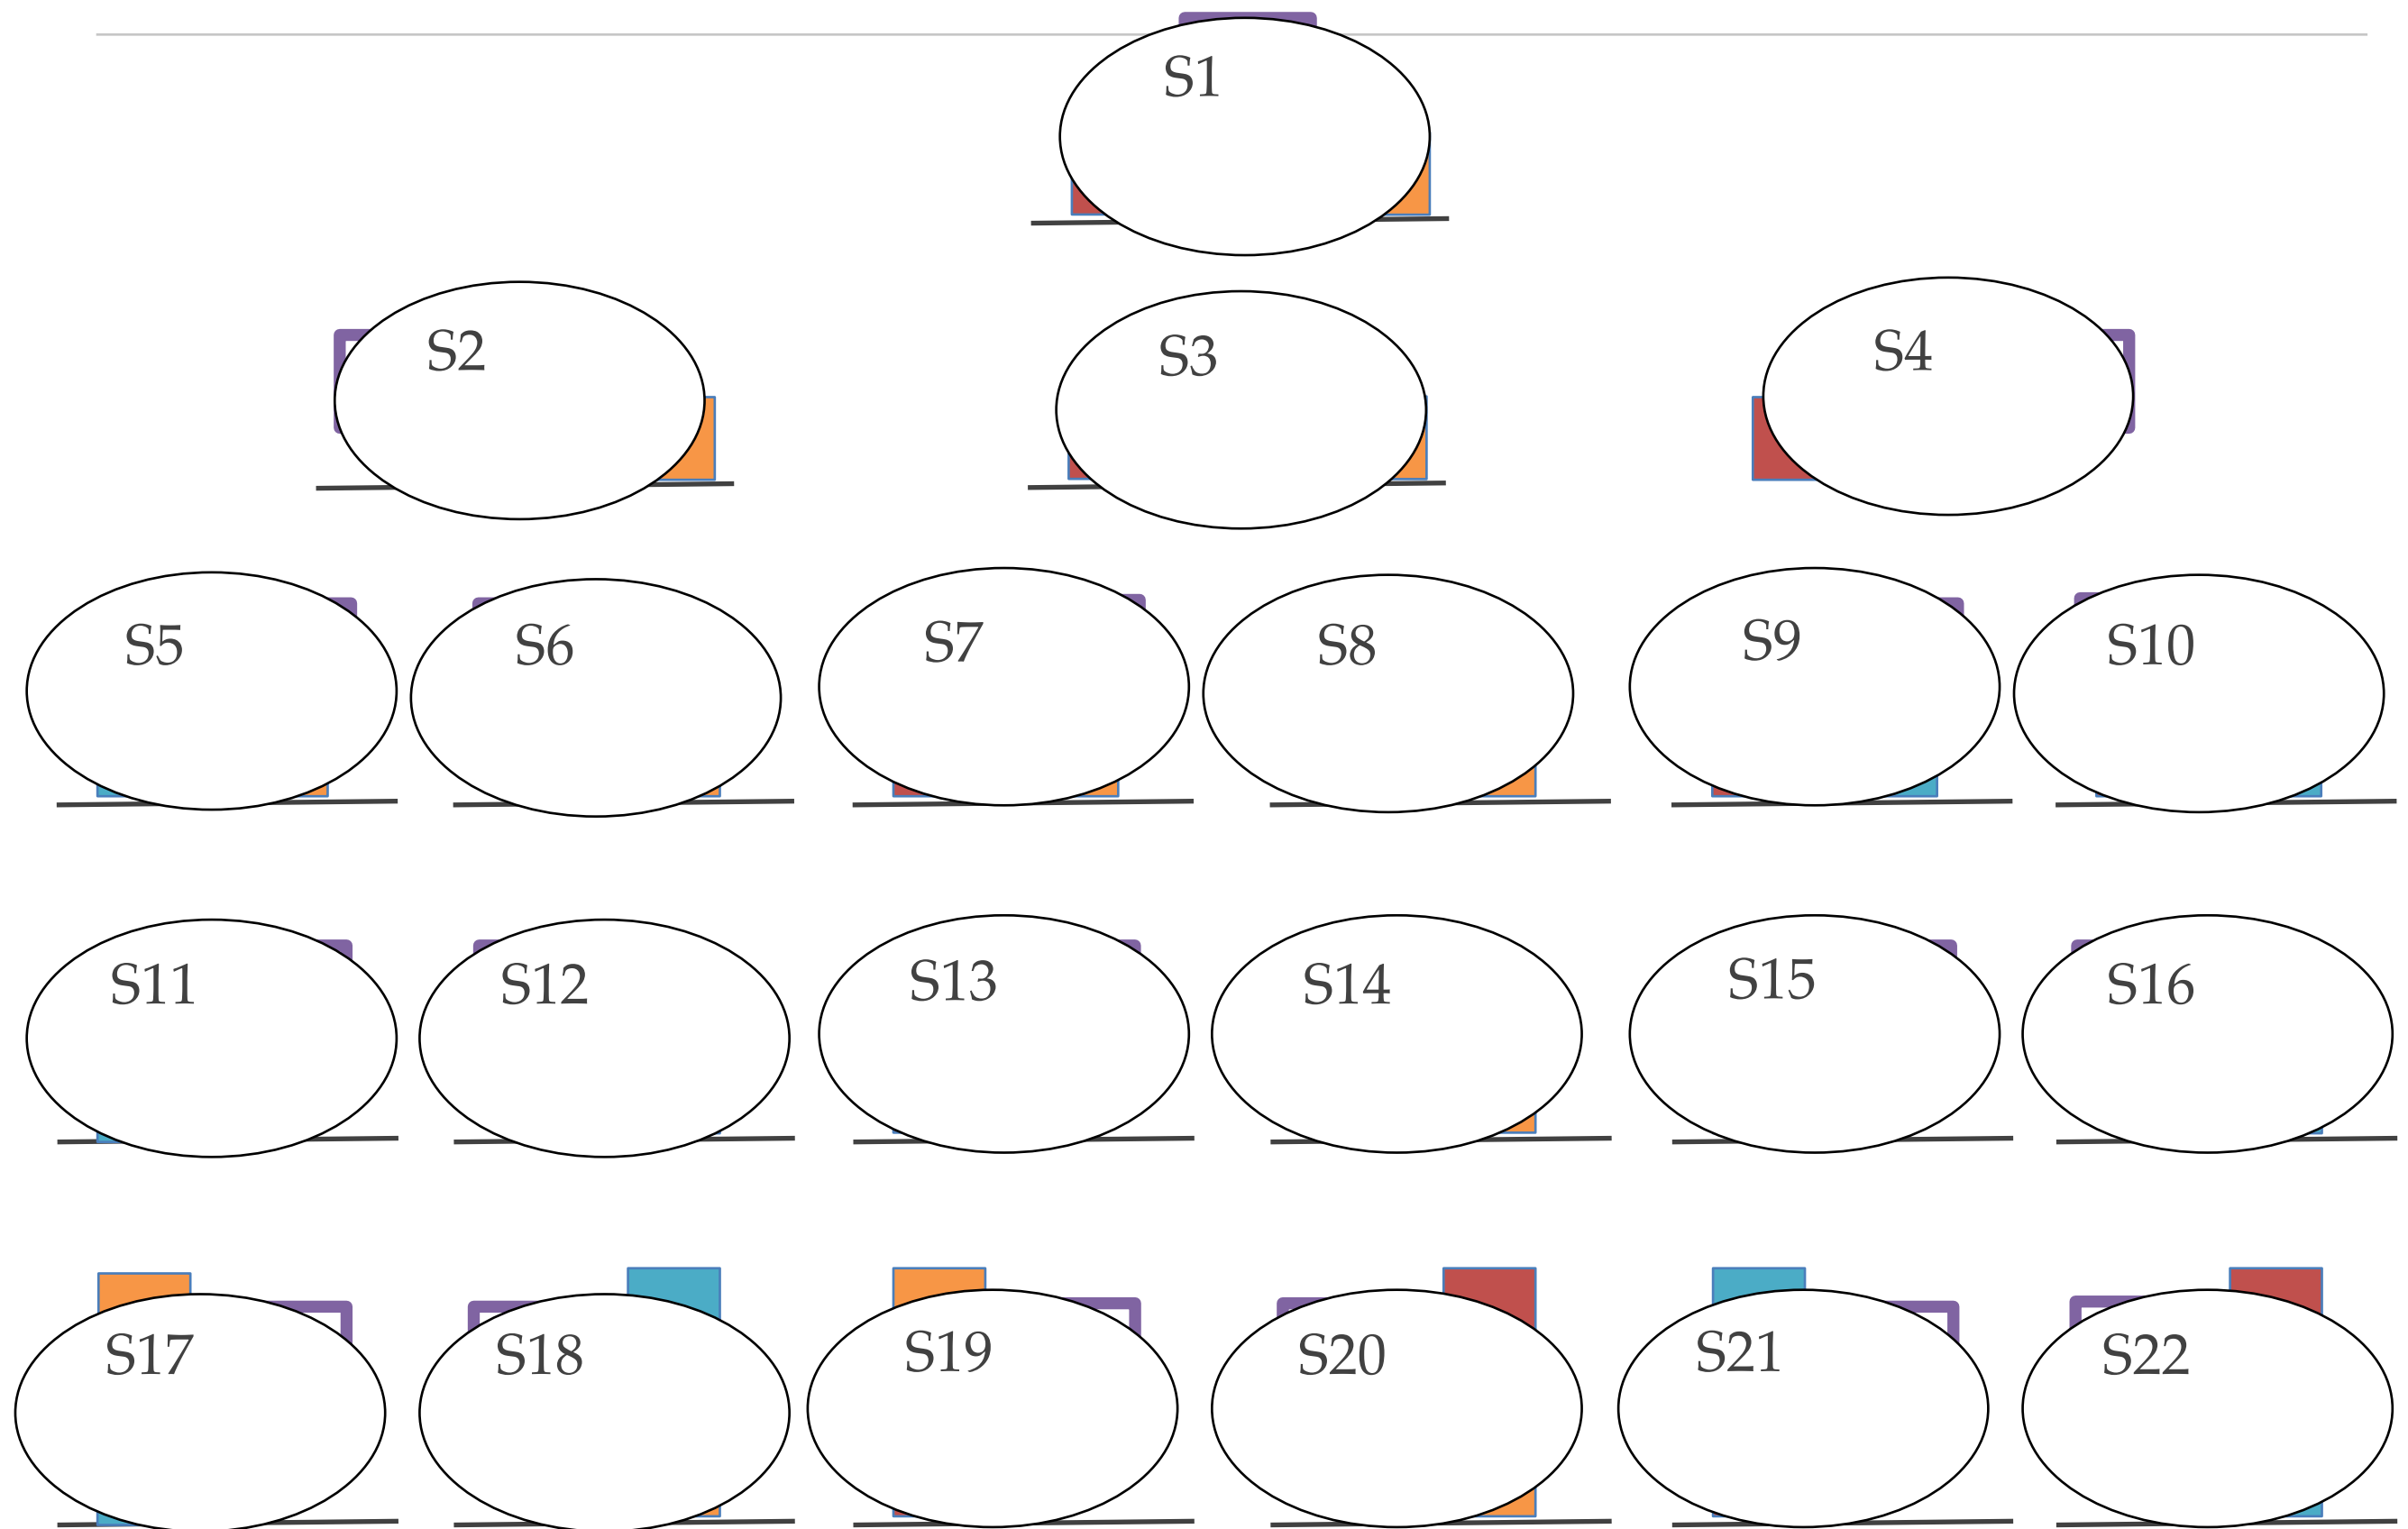
Increasing Generality



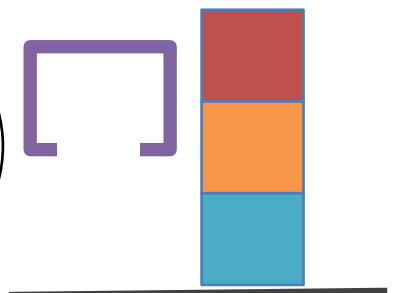
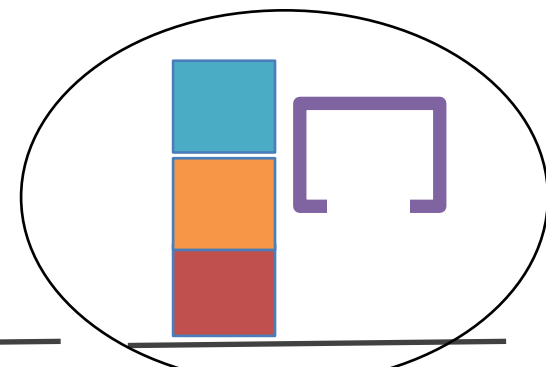

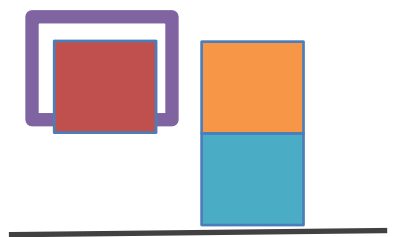
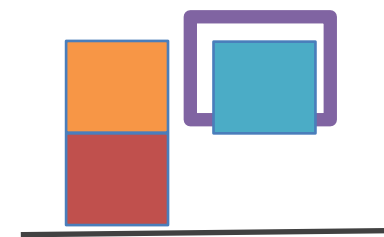

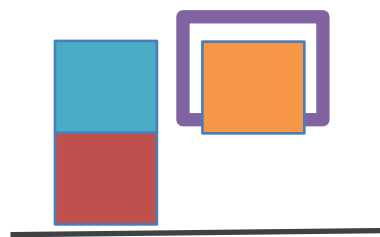
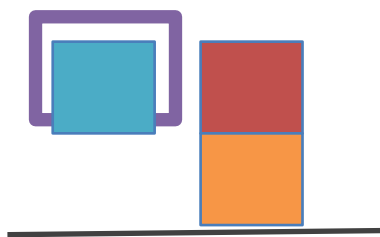

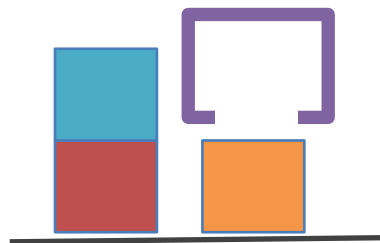



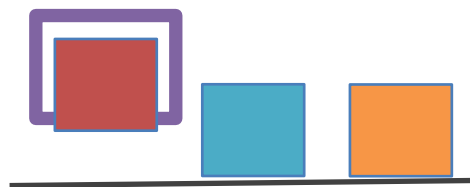
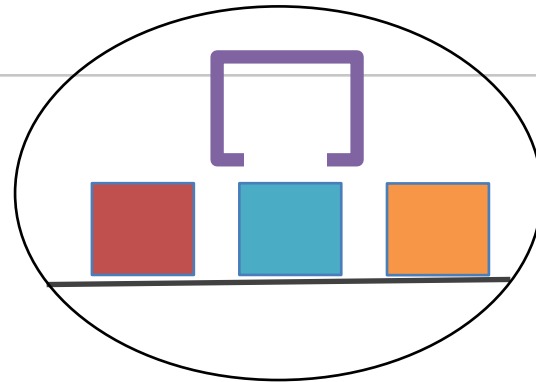
Block Stacking States



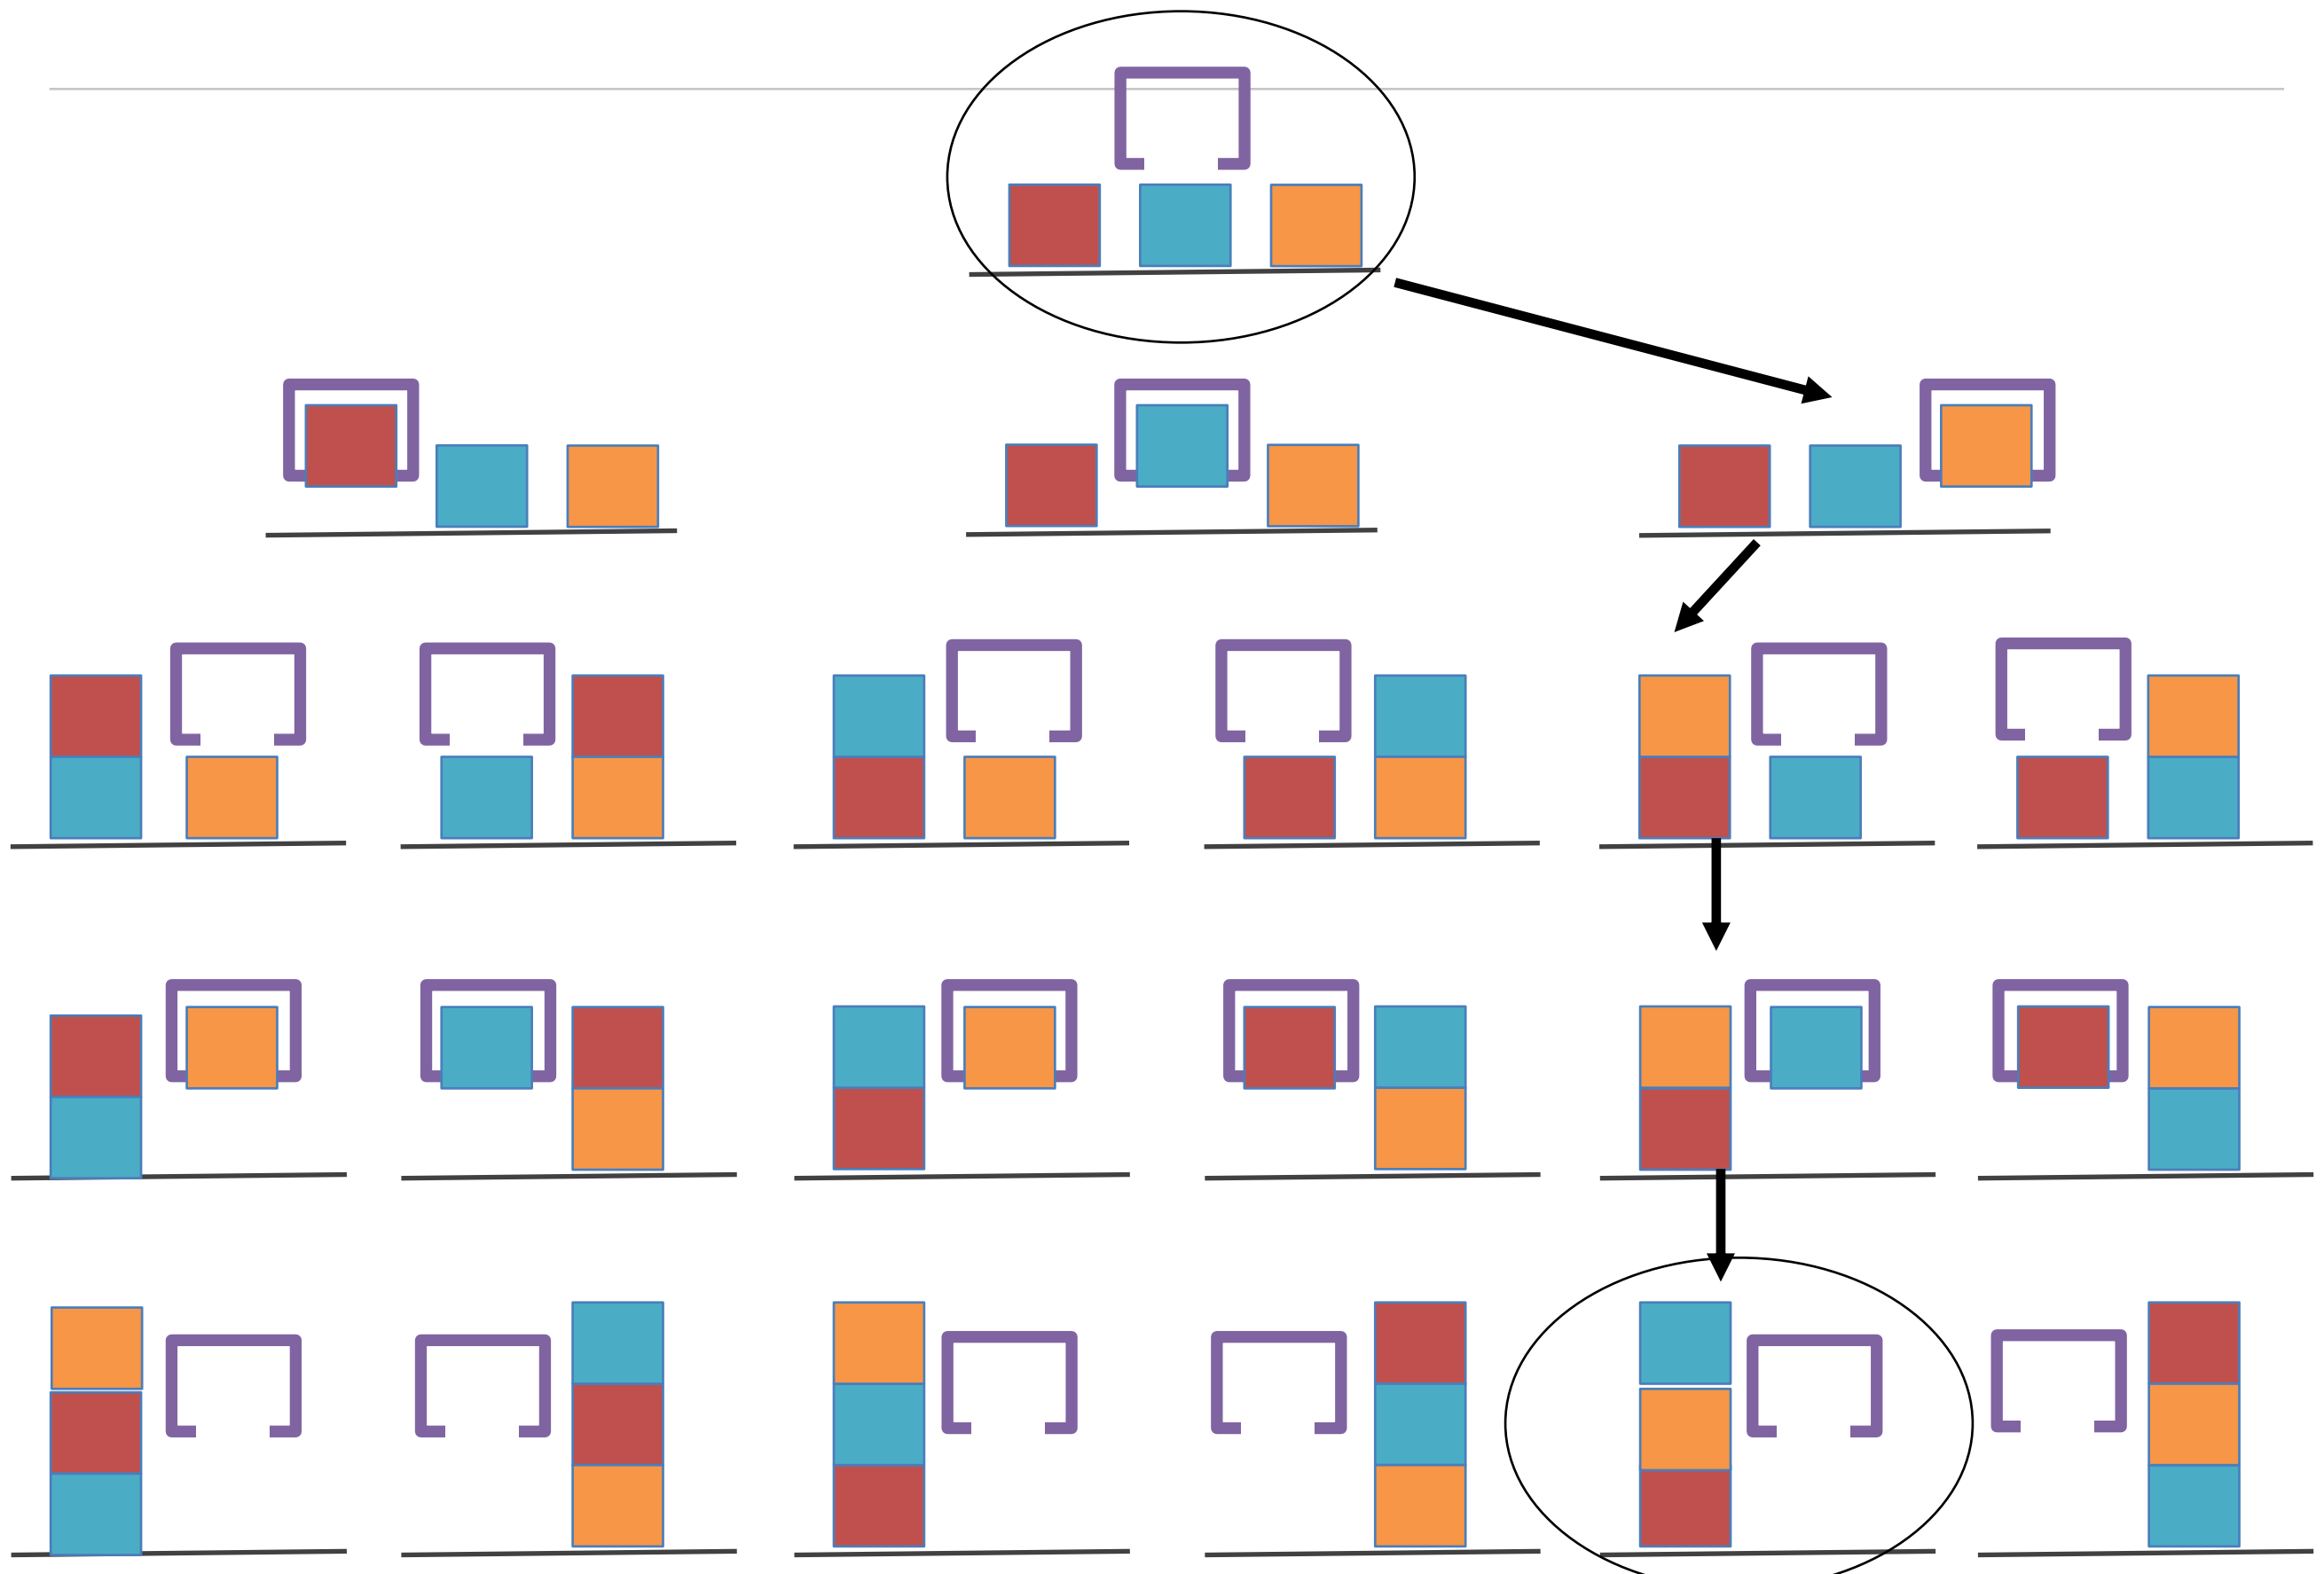
States are Atomic



Initial and Goal States



Plan from Initial to Goal State



Goals in the World

Goal States
completely specified

Goal Statements
partially specified

Preference models
objective function

Increasing Generality



BFS, DFS, A^*

Goals in the World

Goal States
completely specified

Goal Statements
partially specified


Preference models
objective function

Increasing Generality




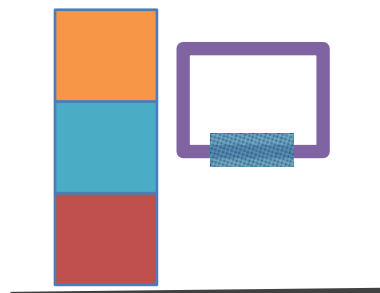
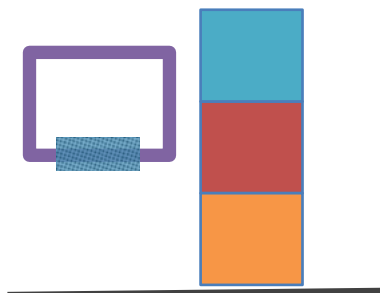
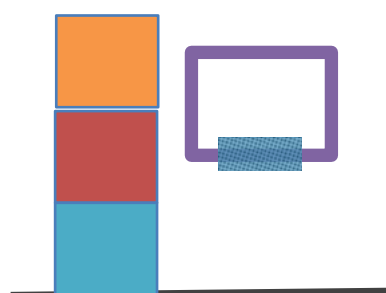



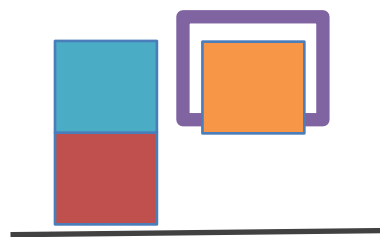
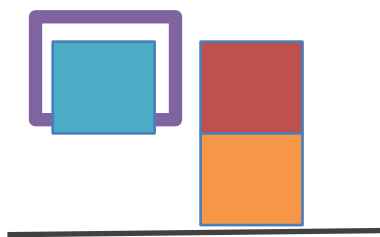
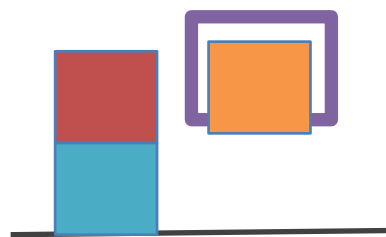



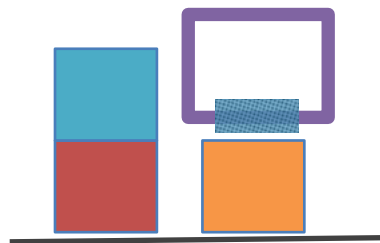
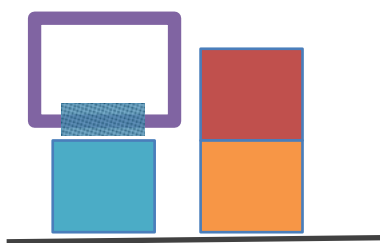
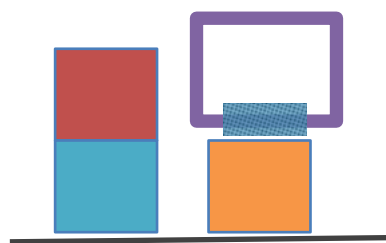


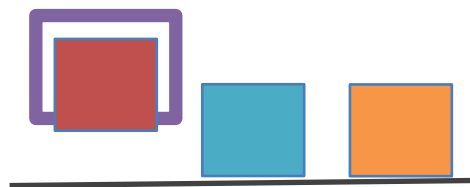


BFS, DFS, A^*

Reward Function



Living cost = -1
Large positive reward in goal
state



Goals in the World

Goal States
completely specified

Goal Statements
partially specified

Preference models
objective function

Increasing Generality



BFS, DFS, A^*

A^* , MDP, RL

Goals in the World

Goal States
completely specified

Goal Statements
partially specified

Preference models
objective function

Increasing Generality



BFS, DFS, A^*

?

A^* , MDP, RL

Goals in the World

Goal States
completely specified

Goal Statements
partially specified

Preference models
objective function

Increasing Generality



BFS, DFS, A^*

Logic, CSP

A^* , MDP, RL

Logical Agents

Create a Knowledge Base (KB)

Symbols – each is true or false

TELL KB

Initial state and a priori knowledge


Domain knowledge and “Physics” of the domain

Logical Agents with PL

Create a Knowledge Base Symbols

A-on-Table	A-In-Hand
B-on-Table	B-In-Hand
C-on-Table	C-In-Hand
Hand-Empty	
A-on-B	B-on-A
A-on-C	C-on-A
B-on-C	C-on-B

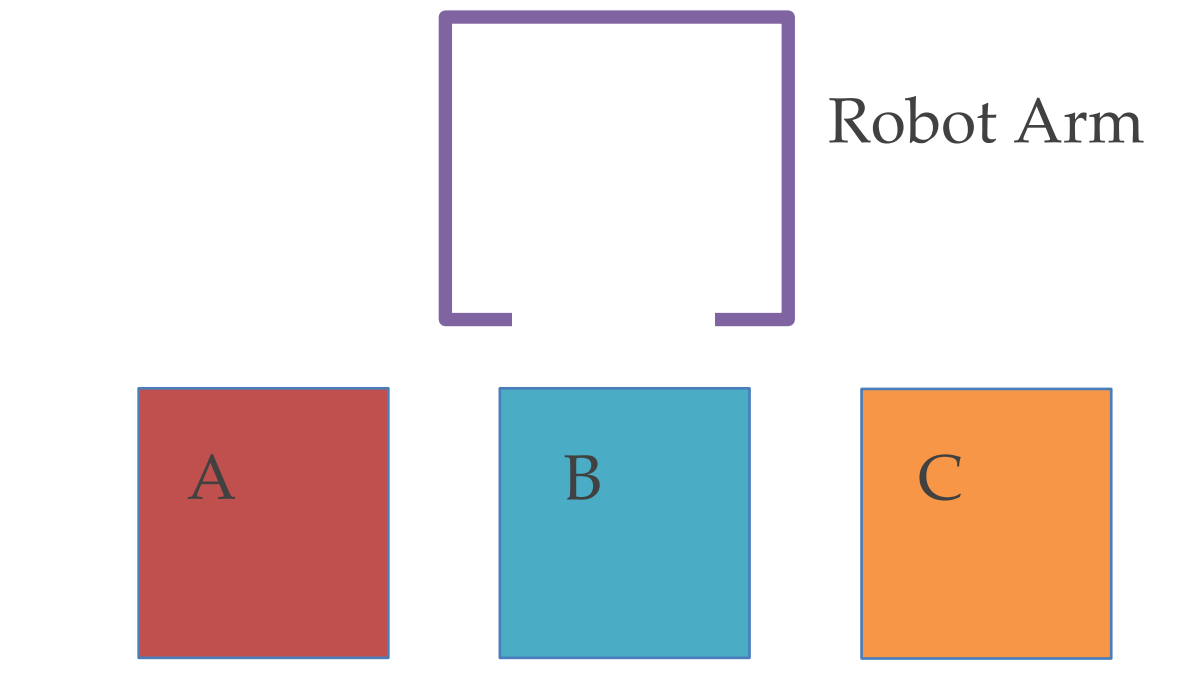
Pick_A, Pick_B, Pick_C
Put_on_Table_A, Put_on_Table_B, Put_on_Table_C
Anything missing?



Logical Agents with PL

Create a Knowledge Base Symbols


A-on-Table[t] A-In-Hand[t]
B-on-Table[t] B-In-Hand[t]
C-on-Table[t] C-In-Hand[t]
Hand-Empty[t]
A-on-B[t] B-on-A[t]
A-on-C[t] C-on-A[t]
B-on-C[t] C-on-B[t]
Pick_A[t], Pick_B[t], Pick_C[t]
Put_on_Table_A[t], Put_on_Table_B[t], Put_on_Table_C[t]
...



Logical Agents with PL

Create a Knowledge Base Symbols

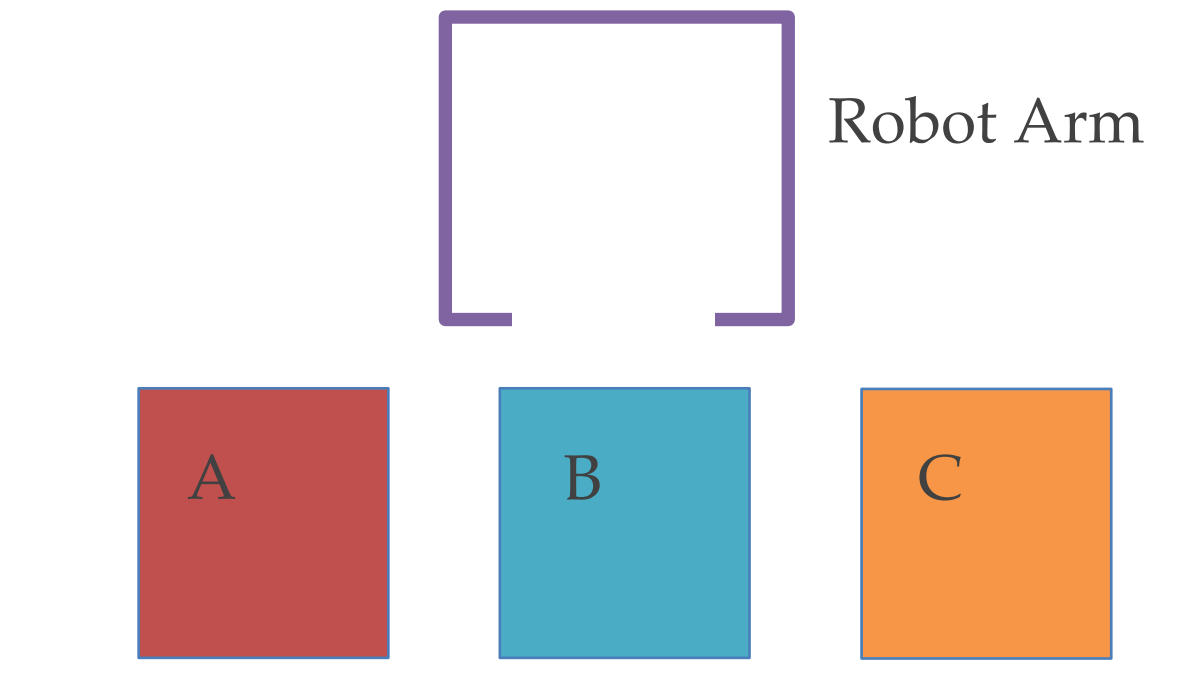
$A\text{-on-Table}[0]$ $A\text{-In-Hand}[t]$
 $B\text{-on-Table}[0]$ $B\text{-In-Hand}[t]$
 $C\text{-on-Table}[0]$ $C\text{-In-Hand}[t]$
 $\text{Hand-Empty}[0]$
 $A\text{-on-}B[t]$ $B\text{-on-}A[t]$
 $A\text{-on-}C[t]$ $C\text{-on-}A[t]$
 $B\text{-on-}C[t]$ $C\text{-on-}B[t]$
 $\text{Pick_}A[t], \text{Pick_}B[t], \text{Pick_}C[t]$
 $\text{Put_on_Table_}A[t], \text{Put_on_Table_}B[t], \text{Put_on_Table_}C[t]$
...



Logical Agents with PL

Create a Knowledge Base
Implications

When is A-on-Table[t] true or false?
Use successor-state axioms!



Logical Agents

Create a Knowledge Base
Implications


$A\text{-on-Table}[t]$

\Leftrightarrow

$(A\text{-on-Table}[t-1] \wedge \neg \text{Pick_A}[t-1])$

\vee

$(A\text{-in-Hand}[t-1] \wedge \text{Put-on-Table}[t-1])$



Logical Agents

Create a Knowledge Base (KB)

Symbols – each is true or false

TELL KB

Initial state and a priori knowledge

Domain knowledge and “Physics” of the domain

ASK whether KB entails a query

e.g., $B\text{-on-}C[t] \wedge C\text{-on-}A[t]$?

Partially-Specified Goal

We didn't specify what all goal symbols needed to be, only some

There are potentially many possible world (i.e., goal states) that could satisfy this goal

We only need to search for one satisfying assignment of variables

Challenges of Logic Planning

We need symbols for each time step (even with FOL)

A state (i.e., model or possible world) is fully-specified by the list of all literals that are true in this model

Actions (e.g., picking a block) are represented with successor-state axioms

Easy to incorrectly specify an axiom (e.g., Hand-Empty[t])

So many symbols means it is hard to debug

Classical Planning (with STRIPS)

Also partially-specified goals

Create a Knowledge Base

- Using STRIPS language

 - Predicates for describing states

 - Operators for describing actions

Using KB, find plan to reach any goal state

- Goal = conjunction of predicates

Predicates

Propositional Logic

A-on-Table, A-In-Hand,
B-on-Table, B-In-Hand,
C-on-Table, C-In-Hand,
Hand-Empty,
A-on-B, B-on-A,
A-on-C, C-on-A,
B-on-C, C-on-B
Clear-A, Clear-B,
Clear-C
...

STRIPS

Constants: A, B, C

Predicates:

In-Hand(A)

On-Table(B)

On-Block(B,C)

HandEmpty()

Clear(A)

...

STRIPS


- ❖ STRIPS inspired by FOL
- ❖ In STRIPS
 - ❖ No functions!
 - ❖ States are represented by conjunctions of positive predicates
 - ❖ If a predicate doesn't appear in this conjunction, it is false
 - ❖ Predicates don't depend on time
 - ❖ Actions don't need any predicates, they are represented as operators
- ❖ Trade-off between expressivity of language and efficiency of solving algorithms

How to Describe this State?

Instances: A, B, C

Propositions:

- 1) In-Hand(A)
- 2) In-Hand(B)
- 3) In-Hand(C)
- 4) On-Table(A)
- 5) On-Table(B)
- 6) On-Table(C)
- 7) On-Block(B,C)
- 8) On-Block(A,B)
- 9) HandEmpty()




Quiz: Check all that apply

Instances: A, B, C

Propositions:

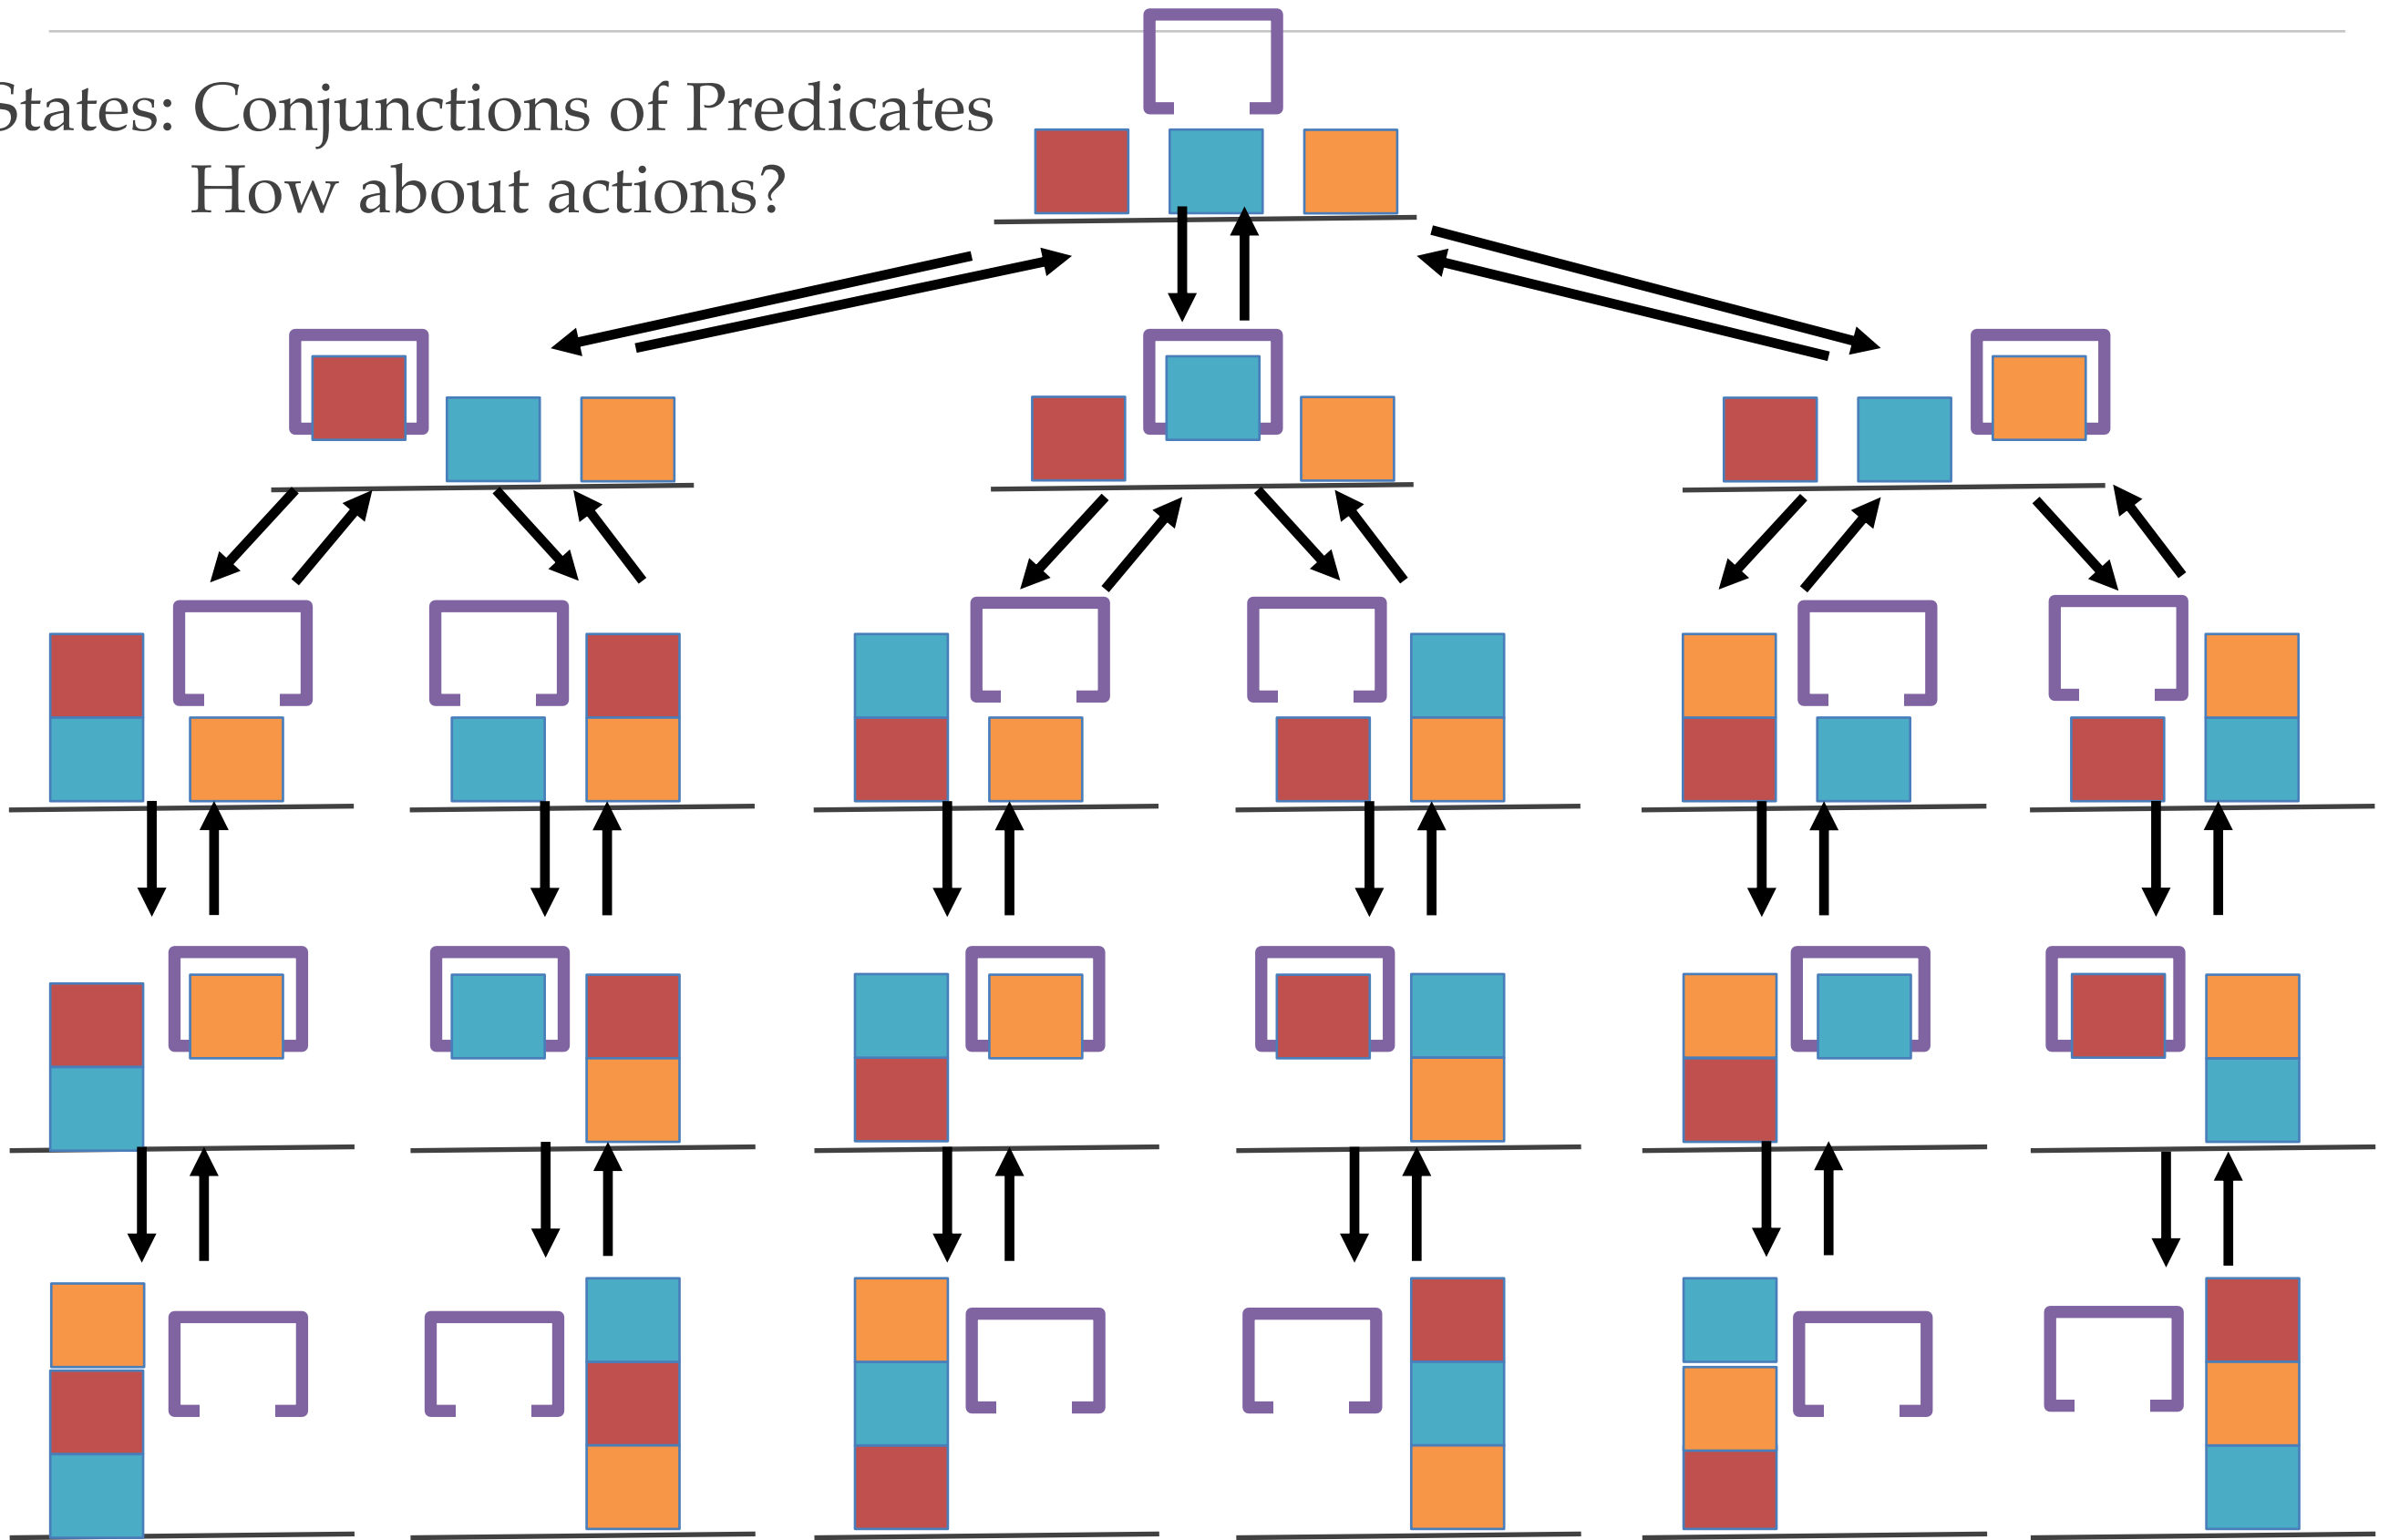
- 1) In-Hand(A)
- 2) In-Hand(B)
- 3) In-Hand(C)
- 4) On-Table(A)
- 5) On-Table(B)
- 6) On-Table(C)
- 7) On-Block(B,C)
- 8) On-Block(A,B)
- 9) HandEmpty()



Block Stacking

States: Conjunctions of Predicates

How about actions?



Operators

Actions can be applied only if some conditions are met

Represented as conjunctions of positive predicates

Actions change the state of the world

Represented as effects that add / delete predicates

Operator = precondition, delete list, add list

Operators can have parameters and are given names

e.g., pick-up(o), put-down(o)

Actions for Block Stacking




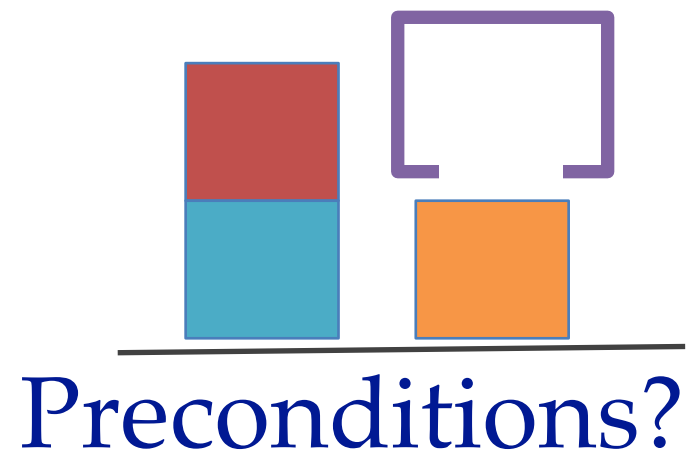
Blocks are picked up and put down by the hand

Blocks can be picked up only if they are clear

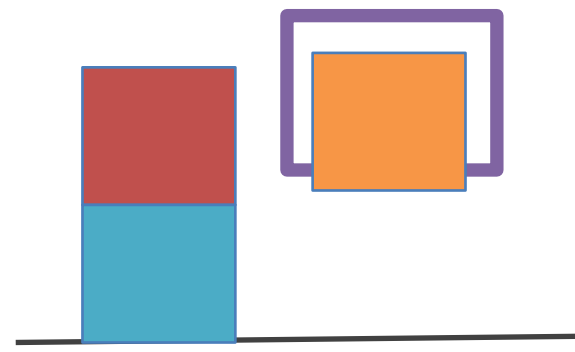
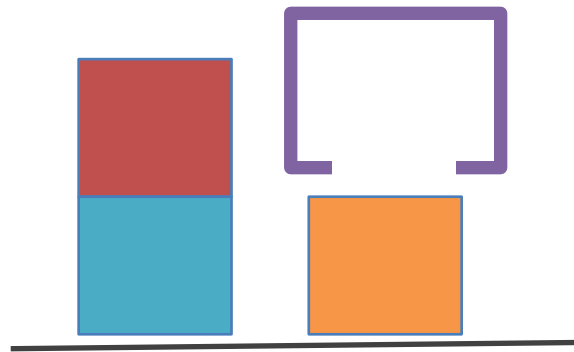
Hand can pick up a block only if the hand is empty

Hand can put down blocks on blocks or on the table

Pick Up Block from Table Example



Pick Up Block from Table Example



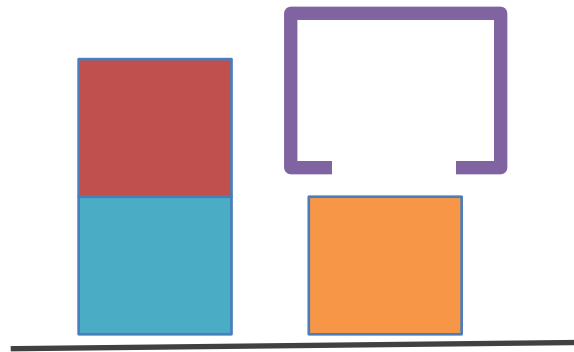
Preconditions

HandEmpty

On-Table(b)

Clear(b)

Pick Up Block from Table Example




Preconditions

HandEmpty

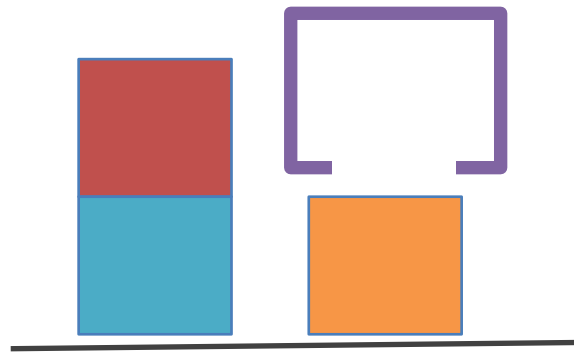
On-Table(b)

Clear(b)



Effects?

Pick Up Block from Table Example

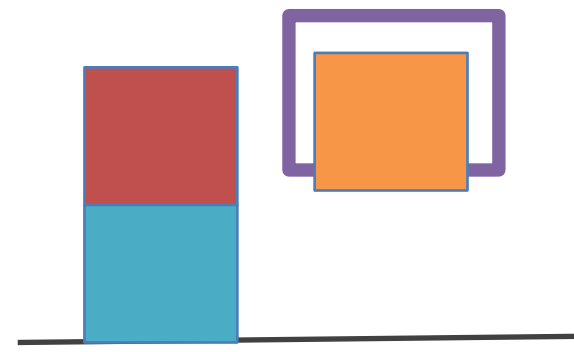


Preconditions

HandEmpty

On-Table(b)

Clear(b)



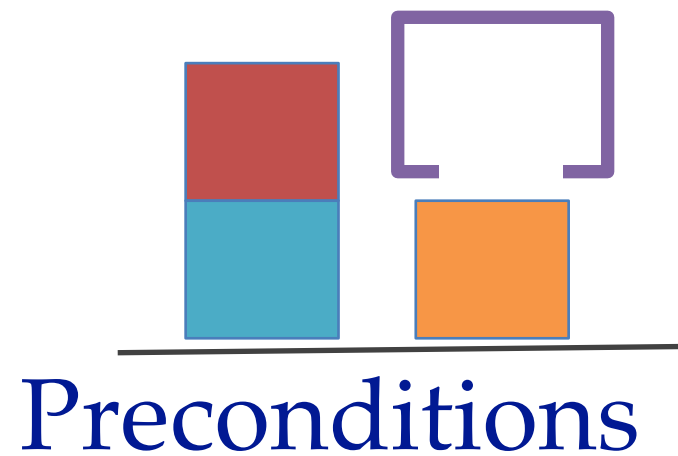
Effects

Add: Holding(b)

Delete: On-Table(b)

HandEmpty

Pick Block from Block Example



Operators for Block Stacking

Pickup_from_Table(b):

Pre: HandEmpty, Clear(b),
On-Table(b)

Add: Holding(b)

Delete: HandEmpty, On-Table(b)

Pickup_from_Block(b,c):

Pre: HandEmpty, On(b,c), $b \neq c$

Add: Holding(b), Clear(c)

Delete: HandEmpty, On(b,c)

Operators for Block Stacking

Pickup_from_Table(b):

Pre: HandEmpty, Clear(b),
On-Table(b)

Add: Holding(b)

Delete: HandEmpty, On-Table(b)

Pickup_from_Block(b,c):

Pre: HandEmpty, On(b,c)

Add: Holding(b), Clear(c)

Delete: HandEmpty, On(b,c)

Putdown_on_Table(b):

Pre: Holding(b)

Add: HandEmpty, On-Table(b)

Delete: Holding(b)

Putdown_on_Block(b,c):

Pre: Holding(b), Clear(c)

Add: HandEmpty, On(b,c)

Delete: Clear(c), Holding(b)

Example Plan of Actions

$\text{HandEmpty} \wedge \text{On-Table(R)} \wedge \text{On(T,R)} \wedge \text{Clear(T)} \wedge \text{On-Table(O)} \wedge \text{Clear(O)}$



Example Plan of Actions

$\text{HandEmpty} \wedge \text{On-Table}(R) \wedge \text{On}(T,R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O)$

Pickup_from_Block(b,c):

Pre: $\text{HandEmpty}, \text{On}(b,c), b \neq c$

Add: $\text{Holding}(b), \text{Clear}(c)$

Delete: $\text{HandEmpty}, \text{On}(b,c)$

Pickup_from_Table(b):

Pre: $\text{HandEmpty}, \text{Clear}(b), \text{On-Table}(b)$

Add: $\text{Holding}(b)$

Delete: $\text{HandEmpty}, \text{On-Table}(b)$



Example Plan of Actions

$\text{HandEmpty} \wedge \text{On-Table}(R) \wedge \text{On}(T,R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O)$
Pickup_from_Block(T,R)

Pickup_from_Block(b,c):
Pre: $\text{HandEmpty}, \text{On}(b,c), \text{Clear}(c), b \neq c$
Add: $\text{Holding}(b), \text{Clear}(c)$
Delete: $\text{HandEmpty}, \text{On}(b,c)$



Example Plan of Actions

$\text{HandEmpty} \wedge \text{On-Table}(R) \wedge \text{On}(T,R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O)$

Pickup_from_Block(T,R)

$\text{On-Table}(R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O)$

Pickup_from_Block(b,c):

Pre: HandEmpty, On(b,c), Clear(c), $b \neq c$

Add: Holding(b), Clear(c)

Delete: HandEmpty, On(b,c)



Example Plan of Actions

$\text{HandEmpty} \wedge \text{On-Table}(R) \wedge \text{On}(T,R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O)$

Pickup_from_Block(T,R)

$\text{On-Table}(R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O) \wedge \text{Holding}(T) \wedge \text{Clear}(R)$

Pickup_from_Block(b,c):
Pre: $\text{HandEmpty}, \text{On}(b,c), \text{Clear}(c), b \neq c$
Add: $\text{Holding}(b), \text{Clear}(c)$
Delete: $\text{HandEmpty}, \text{On}(b,c)$



Example Plan of Actions

$\text{HandEmpty} \wedge \text{On-Table}(R) \wedge \text{On}(T,R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O)$

Pickup_from_Block(T,R)

$\text{On-Table}(R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O) \wedge \text{Holding}(T) \wedge \text{Clear}(R)$



Example Plan of Actions

$\text{HandEmpty} \wedge \text{On-Table}(R) \wedge \text{On}(T,R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O)$

Pickup_from_Block(T,R)

$\text{On-Table}(R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O) \wedge \text{Holding}(T) \wedge \text{Clear}(R)$

Putdown_on_Table(T)

Putdown_on_Table(b):

Pre: Holding(b)

Add: HandEmpty, On-Table(b)

Delete: Holding(b)



Example Plan of Actions

$\text{HandEmpty} \wedge \text{On-Table}(R) \wedge \text{On}(T,R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O)$

Pickup_from_Block(T,R)

$\text{On-Table}(R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O) \wedge \text{Holding}(T) \wedge \text{Clear}(R)$

Putdown_on_Table(T)

$\text{On-Table}(R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O) \wedge \text{Clear}(R)$

Putdown_on_Table(b):

Pre: Holding(b)

Add: HandEmpty, On-Table(b)

Delete: Holding(b)



Example Plan of Actions

$\text{HandEmpty} \wedge \text{On-Table}(R) \wedge \text{On}(T,R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O)$

Pickup_from_Block(T,R)

$\text{On-Table}(R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O) \wedge \text{Holding}(T) \wedge \text{Clear}(R)$

Putdown_on_Table(T)

$\text{On-Table}(R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O) \wedge \text{Clear}(R) \wedge \text{HandEmpty} \wedge \text{On-Table}(T)$

Putdown_on_Table(b):

Pre: Holding(b)

Add: HandEmpty, On-Table(b)

Delete: Holding(b)



Example Plan of Actions

$\text{HandEmpty} \wedge \text{On-Table}(R) \wedge \text{On}(T,R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O)$

Pickup_from_Block(T,R)

$\text{On-Table}(R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O) \wedge \text{Holding}(T) \wedge \text{Clear}(R)$

Putdown_on_Table(T)

$\text{On-Table}(R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O) \wedge \text{Clear}(R) \wedge \text{HandEmpty} \wedge \text{On-Table}(T)$



Example Plan of Actions

$\text{HandEmpty} \wedge \text{On-Table}(R) \wedge \text{On}(T,R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O)$

Pickup_from_Block(T,R)

$\text{On-Table}(R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O) \wedge \text{Holding}(T) \wedge \text{Clear}(R)$

Putdown_on_Table(T)

$\text{On-Table}(R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O) \wedge \text{Clear}(R) \wedge \text{HandEmpty} \wedge \text{On-Table}(T)$

Pickup_from_Table(O)

Pickup_from_Table(b):
Pre: HandEmpty, Clear(b), On-Table(b)
Add: Holding(b)
Delete: HandEmpty, On-Table(b)



Example Plan of Actions

$\text{HandEmpty} \wedge \text{On-Table}(R) \wedge \text{On}(T,R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O)$

Pickup_from_Block(T,R)

$\text{On-Table}(R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O) \wedge \text{Holding}(T) \wedge \text{Clear}(R)$

Putdown_on_Table(T)

$\text{On-Table}(R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O) \wedge \text{Clear}(R) \wedge \text{HandEmpty} \wedge \text{On-Table}(T)$

Pickup_from_Table(O)

$\text{On-Table}(R) \wedge \text{Clear}(T) \wedge \text{Clear}(O) \wedge \text{Clear}(R) \wedge \text{On-Table}(T) \wedge \text{Holding}(O)$

Pickup_from_Table(b):

Pre: HandEmpty, Clear(b), On-Table(b)

Add: Holding(b)

Delete: HandEmpty, On-Table(b)



Example Plan of Actions

$\text{HandEmpty} \wedge \text{On-Table}(R) \wedge \text{On}(T,R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O)$

Pickup_from_Block(T,R)

$\text{On-Table}(R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O) \wedge \text{Holding}(T) \wedge \text{Clear}(R)$

Putdown_on_Table(T)

$\text{On-Table}(R) \wedge \text{Clear}(T) \wedge \text{On-Table}(O) \wedge \text{Clear}(O) \wedge \text{Clear}(R) \wedge \text{HandEmpty} \wedge \text{On-Table}(T)$

Pickup_from_Table(O)

$\text{On-Table}(R) \wedge \text{Clear}(T) \wedge \text{Clear}(O) \wedge \text{Clear}(R) \wedge \text{On-Table}(T) \wedge \text{Holding}(O)$

Putdown_on_Block(O,R)

$\text{On-Table}(R) \wedge \text{Clear}(T) \wedge \text{Clear}(O) \wedge \text{On-Table}(T) \wedge \text{On}(O,R) \wedge \text{HandEmpty}$



Planning is a Search Problem

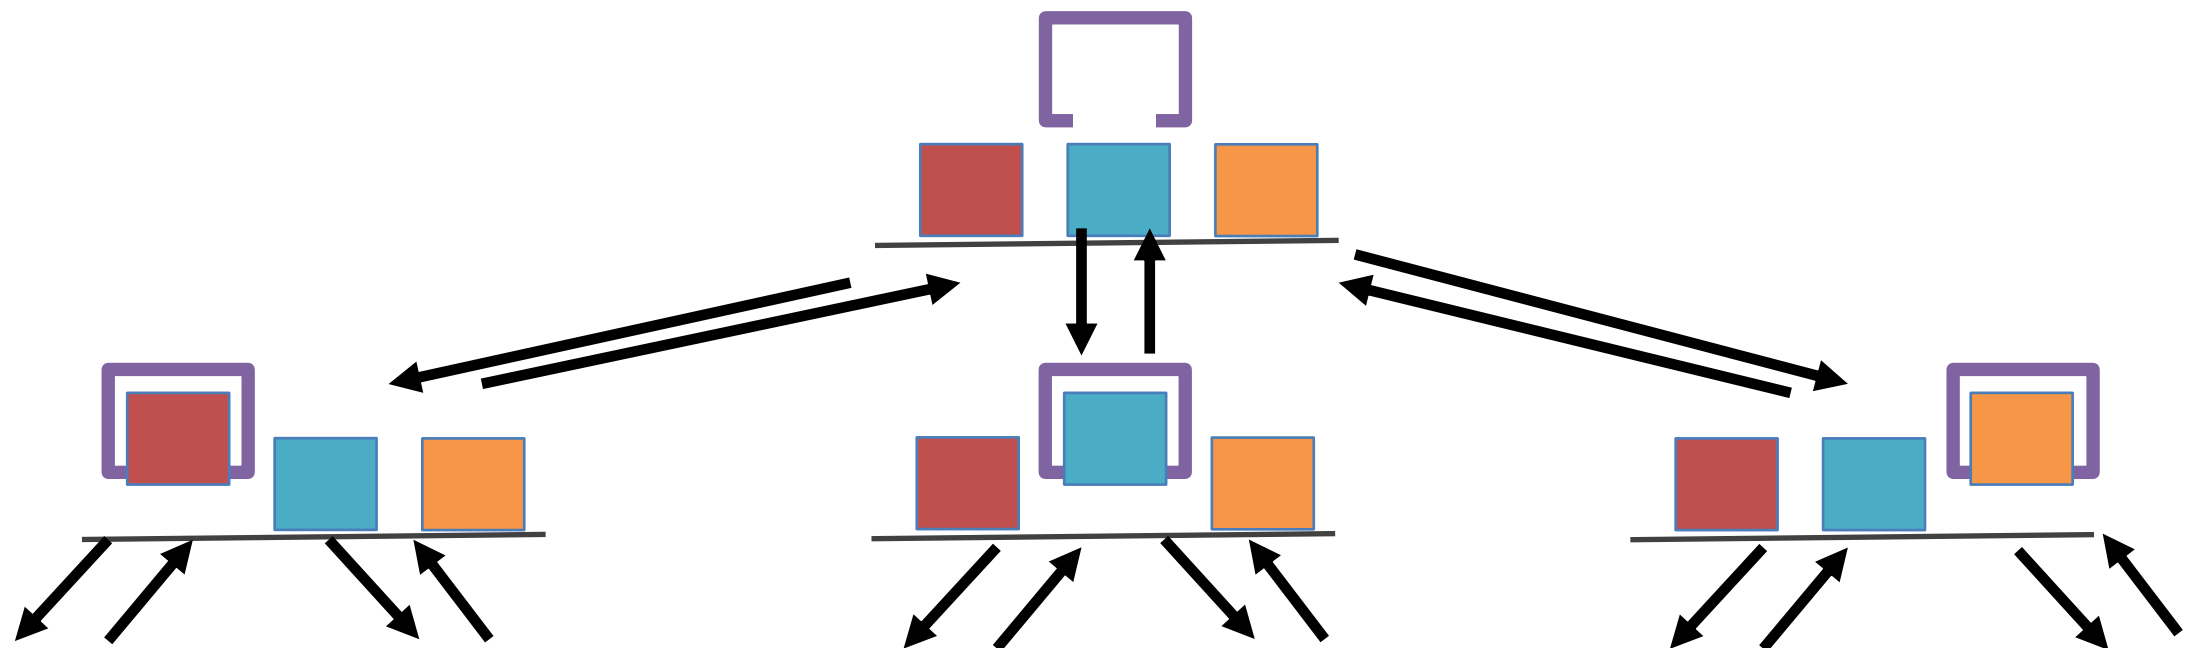
Planning problem can be represented as a graph:

States: Conjunctions of Predicates

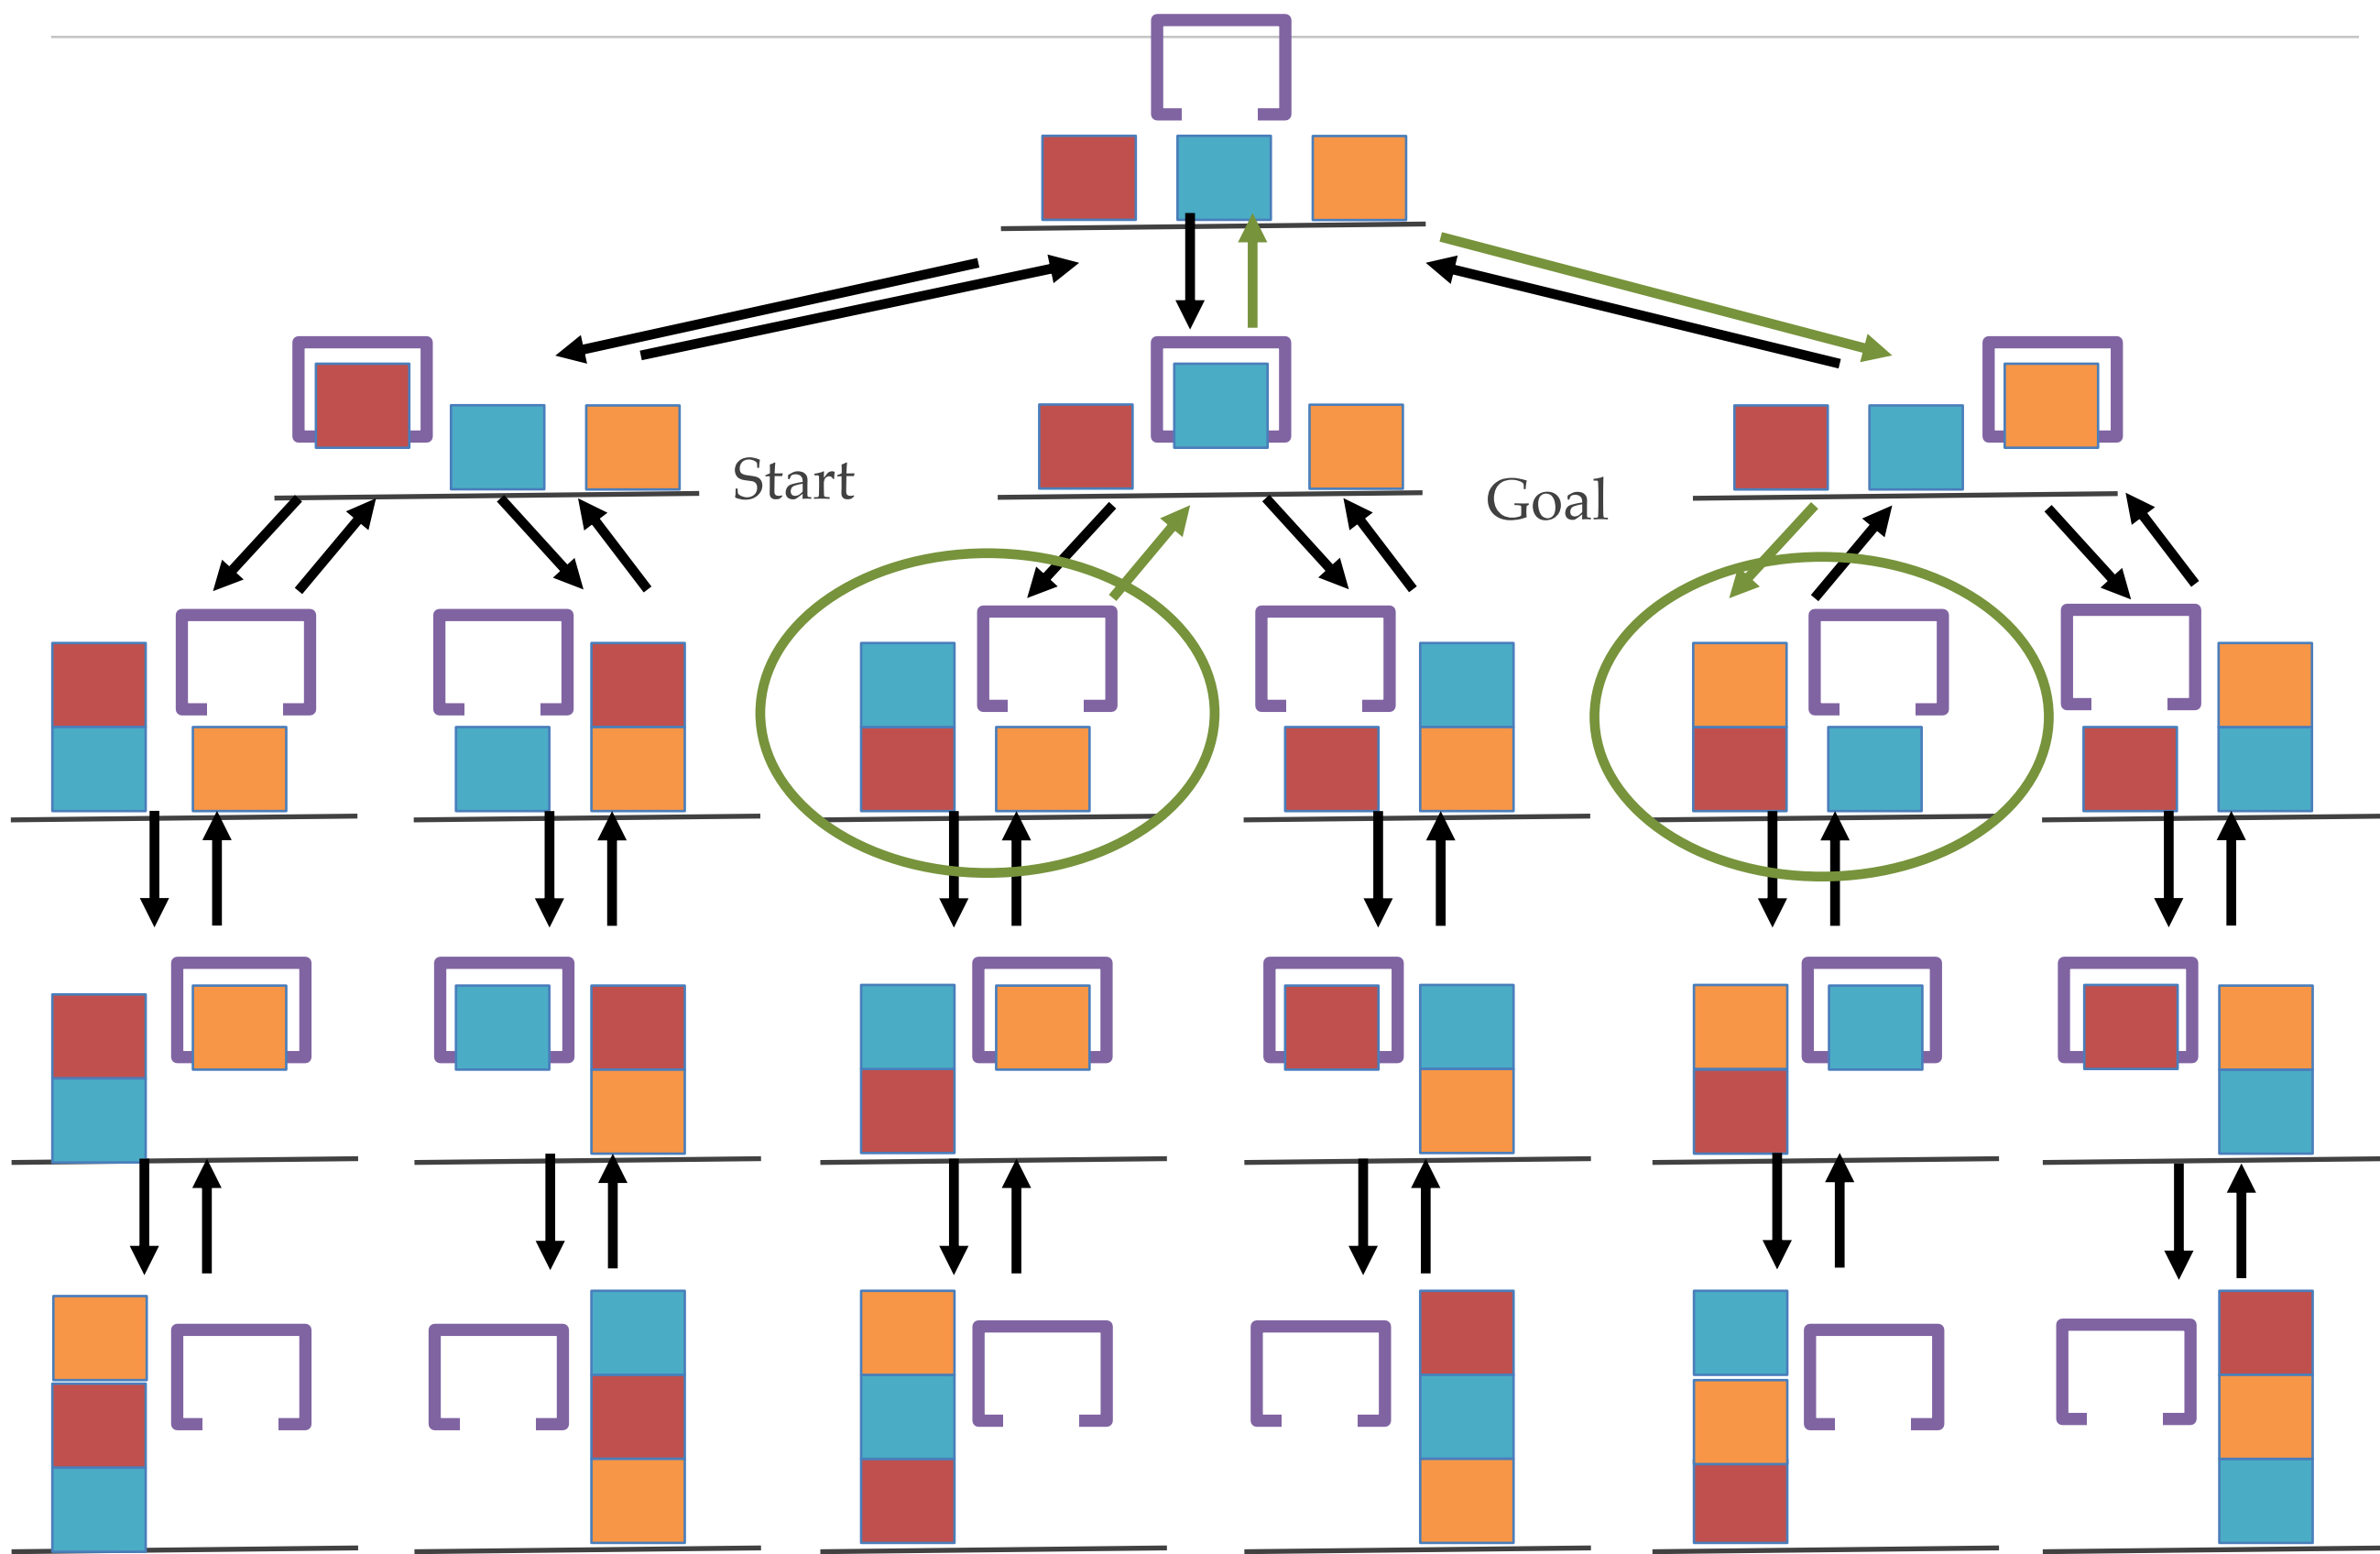
Arrows = Actions

Space complexity in terms of # predicates p ?

Search algorithm can be applied!



Example



Planners

- ❖ Any search algorithm (BFS, DFS, IDS...) could be used
- ❖ Here, search can also be performed backward!
 - ❖ From a given state $X_1 \wedge \dots \wedge X_k$, consider relevant and consistent actions
 - ❖ An action is relevant if it achieves one of the X_i 's
 - ❖ An action is consistent if it doesn't undo one of the X_i 's
 - ❖ An operator can easily be reversed by modifying the current state:
 - ❖ Remove its positive effects
 - ❖ Add its preconditions (unless it already appears in the state)
 - ❖ Stop backward search when reaching state that is satisfied by initial state
- ❖ However, as STRIPS allow to describe compactly very large problems, A^* is needed with very good heuristics

Efficient Algorithms

- ❖ STRIPS planning is PSPACE-complete
- ❖ Forward or backward A*
 - ❖ Domain-independent heuristics
 - ❖ # of unsatisfied literals in goal
 - ❖ Relaxed problem
 - ❖ sum of costs for achieving each literal in goal
 - ❖ may be inadmissible
- ❖ Specialized algorithms
 - ❖ e.g., Graphplan (See 10.3 in AIMA)

Properties of Planners

Soundness

- ❖ A planning algorithm is *sound* if all solutions found are legal plans

Completeness

- ❖ A planning algorithm is *complete* if a solution can be found whenever one actually exists

Optimality

- ❖ A planning algorithm is *optimal* if the solution optimizes some measure of plan quality


Future

HOW FRIGHTENED SHOULD WE BE OF A.I.?

Thinking about artificial intelligence can help clarify what makes us human—for better and for worse.



By Tad Friend May 7, 2018



An A.I. system may need to take charge in order to achieve the goals we gave it.

Illustration by Harry Campbell

Combining Symbolic and Connexionist Approaches

Mastering Atari, Go, Chess and Shogi by Planning with a Learned Model

Julian Schrittwieser,^{1*} Ioannis Antonoglou,^{1,2*} Thomas Hubert,^{1*}
Karen Simonyan,¹ Laurent Sifre,¹ Simon Schmitt,¹ Arthur Guez,¹
Edward Lockhart,¹ Demis Hassabis,¹ Thore Graepel,^{1,2} Timothy Lillicrap,¹
David Silver^{1,2*}

¹DeepMind, 6 Pancras Square, London N1C 4AG.

²University College London, Gower Street, London WC1E 6BT.

*These authors contributed equally to this work.

Abstract

Constructing agents with planning capabilities has long been one of the main challenges in the pursuit of artificial intelligence. Tree-based planning methods have enjoyed huge success in challenging domains, such as chess and Go, where a perfect simulator is available. However, in real-world problems the dynamics governing the environment are often complex and unknown. In this work we present the *MuZero* algorithm which, by combining a tree-based search with a learned model, achieves superhuman performance in a range of challenging and visually complex domains, without any knowledge of their underlying dynamics. *MuZero* learns a model that, when applied iteratively, predicts the quantities most directly relevant to planning: the reward, the action-selection policy, and the value function. When evaluated on 57 different Atari games - the canonical video game environment for testing AI techniques, in which model-based planning approaches have historically struggled - our new algorithm achieved a new state of the art. When evaluated on Go, chess and shogi, without any knowledge of the game rules, *MuZero* matched the superhuman performance of the *AlphaZero* algorithm that was supplied with the game rules.

Should we be afraid of AI?

